

# Machine Learning

## Lecture 6: Optimization

---

Prof. Dr. Stephan Günnemann

Data Analytics and Machine Learning  
Technical University of Munich

Winter term 2022/2023

# Motivation

- Many machine learning tasks are optimization problems
- Examples we've already seen:
  - Linear Regression  $\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{2}(\mathbf{X}\mathbf{w} - \mathbf{y})^T(\mathbf{X}\mathbf{w} - \mathbf{y})$  *Ls*
  - Logistic Regression  $\mathbf{w}^* = \arg \min_{\mathbf{w}} -\ln p(\mathbf{y} | \mathbf{w}, \mathbf{X})$  *negative log*
- Other examples:
  - Support Vector Machines: find hyperplane that separates the classes with a maximum margin
  - k-means: find clusters and centroids such that the squared distances is minimized
  - Matrix Factorization: find matrices that minimize the reconstruction error
  - Neural networks: find weights such that the loss is minimized
  - And many more...

# General Task


- Let  $\theta$  denote the variables/parameters of our problem we want to learn
  - e.g.  $\theta = w$  in Logistic Regression
- Let  $\mathcal{X}$  denote the domain of  $\theta$ ; the set of valid instantiations
  - constraints on the parameters!
  - e.g.  $\mathcal{X}$  = set of (positive) real numbers
- Let  $f(\theta)$  denote the **objective function**
  - e.g.  $f$  is the negative log likelihood
- Goal: Find solution  $\theta^*$  minimizing function  $f : \theta^* = \arg \min_{\theta \in \mathcal{X}} f(\theta)$ 
  - find a global minimum of the function  $f$ !
  - similarly, for some problems we are interested in finding the maximum

# Introductory Example

- Goal: Find minimum of function

$$f(\theta) = 0.6 * \theta^4 - 5 * \theta^3 + 13 * \theta^2 - 12 * \theta + 5$$

- Unconstrained optimization + differentiable function
- Necessary condition for minima
  - Gradient = 0
  - Sufficient?





- General challenge: multiple local minima possible

# Convexity: Sets

- $X$  is a convex set

iff

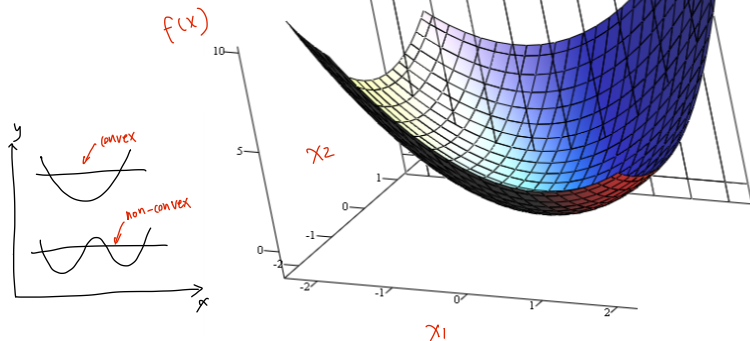

for all  $x, y \in X$  it follows that  $\lambda x + (1 - \lambda)y \in X$  for  $\lambda \in [0, 1]$



# Convexity: Functions

- $f(x)$  is a convex function on the convex set  $X$  iff

for all  $x, y \in X : \lambda f(x) + (1 - \lambda)f(y) \geq f(\lambda x + (1 - \lambda)y)$  for  $\lambda \in [0, 1]$



# Convexity and *minimization problems*

- Region **above** a convex function is convex



$$f(\lambda \mathbf{x} + (1 - \lambda) \mathbf{y}) \leq \lambda f(\mathbf{x}) + (1 - \lambda) f(\mathbf{y})$$

hence  $\lambda f(\mathbf{x}) + (1 - \lambda) f(\mathbf{y}) \in X$  for  $\mathbf{x}, \mathbf{y} \in X$

- Convex functions have no local minima which are not global minima
  - Proof by contradiction - linear interpolation breaks local minimum condition



- Each **local minimum** is a **global minimum**
  - zero gradient implies (local) minimum for convex functions
  - if  $f_0$  is a convex function and  $\nabla f_0(\boldsymbol{\theta}^*) = 0$  then  $\boldsymbol{\theta}$  is a global minimum
  - minimization becomes "relatively easy"

# Convexity and *minimization problems*

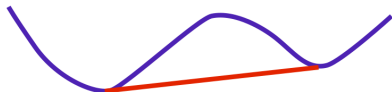
- Region **above** a convex function is convex



$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$$

hence  $\lambda f(x) + (1 - \lambda)f(y) \in X$  for  $x, y \in X$

- Convex functions have no local minima which are not global minima
  - Proof by contradiction - linear interpolation breaks local minimum condition



multiple global minimum

- Each local minimum is a global minimum
  - zero gradient implies (local) minimum for convex functions
  - if  $f_0$  is a convex function and  $\nabla f_0(\theta^*) = 0$  then  $\theta$  is a global minimum
  - minimization becomes "relatively easy"




# First order convexity conditions (I)

- Convexity imposes a rate of rise on the function

- $f((1-t)x + ty) \leq (1-t)f(x) + tf(y)$

- $f(y) - f(x) \geq \frac{f((1-t)x + ty) - f(x)}{t}$

- Difference between  $f(y)$  and  $f(x)$  is bounded by function values between  $x$  and  $y$



# First order convexity conditions (II)

- $f(\mathbf{y}) - f(\mathbf{x}) \geq \frac{f((1-t)\mathbf{x} + t\mathbf{y}) - f(\mathbf{x})}{t}$

- Let  $t \rightarrow 0$  and apply the definition of the derivative

- $f(\mathbf{y}) - f(\mathbf{x}) \geq (\mathbf{y} - \mathbf{x})^T \nabla f(\mathbf{x})$

$$\left[ \frac{f(\mathbf{x} + t(\mathbf{y} - \mathbf{x})) - f(\mathbf{x})}{t} \right] (\mathbf{y} - \mathbf{x})$$

$t \rightarrow 0 \quad \nabla f(\mathbf{x})$

- Theorem:

Suppose  $f : X \rightarrow \mathbb{R}$  is a differentiable function and  $X$  is convex. Then  $f$  is convex iff for  $\mathbf{x}, \mathbf{y} \in X$

$$f(\mathbf{y}) \geq f(\mathbf{x}) + (\mathbf{y} - \mathbf{x})^T \nabla f(\mathbf{x})$$

- Proof. See Boyd p.70

# Verifying convexity (I)

- Convexity makes optimization "easier"
- How to verify whether a function is convex?
- For example:  $e^{x_1+2x_2} + x_1 - \log(x_2)$  convex on  $[1, \infty) \times [1, \infty)$ ?

# Verifying convexity (I)

- Convexity makes optimization "easier"
  - How to verify whether a function is convex?
  - For example:  $e^{x_1+2x_2} + x_1 - \log(x_2)$  convex on  $[1, \infty) \times [1, \infty)$ ?
1. Prove whether the definition of convexity holds (See slide 6)

# Verifying convexity (I)

- Convexity makes optimization "easier"
  - How to verify whether a function is convex?
  - For example:  $e^{x_1+2x_2} + x_1 - \log(x_2)$  convex on  $[1, \infty) \times [1, \infty)$ ?
1. Prove whether the definition of convexity holds (See slide 6)
  2. Exploit special results

# Verifying convexity (I)

- Convexity makes optimization "easier"
  - How to verify whether a function is convex?
  - For example:  $e^{x_1+2x_2} + x_1 - \log(x_2)$  convex on  $[1, \infty) \times [1, \infty)$ ?
1. Prove whether the definition of convexity holds (See slide 6)
  2. Exploit special results
    - First order convexity (See slide 9)
    - Example: A twice differentiable function of one variable is convex on an interval if and only if its second-derivative is non-negative on this interval
    - More general: a twice differentiable function of several variables is convex (on a convex set) if and only if its Hessian matrix is positive semidefinite (on the set)

# Verifying convexity (II)



3. Show that the function can be obtained from simple convex functions by operations that preserve convexity

a) Start with simple convex functions, e.g.

- $f(x) = \text{const}$  and  $f(x) = x^T \cdot b$  (there are also concave functions)
- $f(x) = e^x$

b) Apply "construction rules" (next slide)

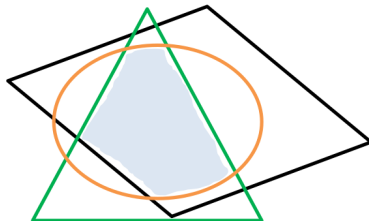
# Convexity preserving operations

- Let  $f_1 : \mathbb{R}^d \rightarrow \mathbb{R}$  and  $f_2 : \mathbb{R}^d \rightarrow \mathbb{R}$  be **convex** functions, and  $g : \mathbb{R}^d \rightarrow \mathbb{R}$  be a **concave** function, then
  - $h(\mathbf{x}) = f_1(\mathbf{x}) + f_2(\mathbf{x})$  is convex
  - $h(\mathbf{x}) = \max\{f_1(\mathbf{x}), f_2(\mathbf{x})\}$  is convex
  - $h(\mathbf{x}) = c \cdot f_1(\mathbf{x})$  is convex if  $c \geq 0$
  - $h(\mathbf{x}) = c \cdot g(\mathbf{x})$  is convex if  $c \leq 0$
  - $h(\mathbf{x}) = f_1(\mathbf{A}\mathbf{x} + \mathbf{b})$  is convex ( $\mathbf{A}$  matrix,  $\mathbf{b}$  vector)
  - $h(\mathbf{x}) = m(f_1(\mathbf{x}))$  is convex if  $m : \mathbb{R} \rightarrow \mathbb{R}$  is convex and nondecreasing
- Example:  $\underbrace{e^{x_1+2x_2} + x_1}_{\text{(convex)}} - \underbrace{\log(x_2)}_{(-) \cdot \text{concave}}$  is convex on, e.g.,  $[1, \infty) \times [1, \infty)$



# Verifying convexity of sets

1. Prove definition
  - often easier for sets than for functions
2. Apply intersection rule
  - Let  $A$  and  $B$  be convex sets, then  $A \cap B$  is a convex set

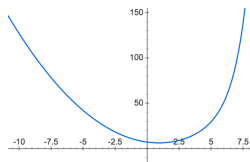


# An easy problem

Convex objective function  $f$

- Objective function differentiable on its whole domain
    - i.e. we are able to compute gradient  $f'$  at every point
  - We can solve  $f'(\theta) = 0$  for  $\theta$  analytically
    - i.e. solution for  $\theta$  where gradient = 0 is known
  - Unconstrained minimization
    - i.e. above computed solution for  $\theta$  is valid
  - We are done!
- 
- Example: Ordinary Least Squares Regression

$$x^2 + e^{x-3} - 2x + 7$$

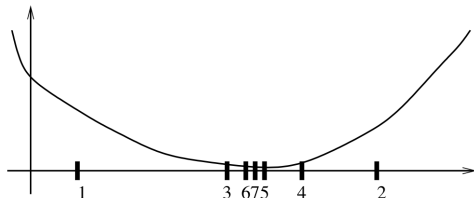


# Outlook

- Unfortunately, many problems are harder...
- No analytical solution for  $f'(\theta) = 0$ 
  - e.g. Logistic Regression
  - Solution: try numerical approaches, e.g. gradient descent
- Constraint on  $\theta$ 
  - e.g.  $f'(\theta) = 0$  only holds for points outside the domain
  - Solution: constrained optimization
- $f$  not differentiable on the whole domain
  - Potential solution: subgradients; or is it a discrete optimization problem?
- $f$  not convex
  - Potential solution: convex relaxations; convex in some variables?

# One-dimensional problems

- Key Idea
  - For differentiable  $f$  search for  $\theta$  with  $\nabla f(\theta) = 0$
  - Interval bisection (derivative is monotonic)



**Require:**  $a, b$ , Precision  $\epsilon$

Set  $A = a, B = b$

**repeat**

**if**  $f'(\frac{A+B}{2}) > 0$  **then**

$$B = \frac{A+B}{2}$$

**else**

$$A = \frac{A+B}{2}$$

**end if**

**until**

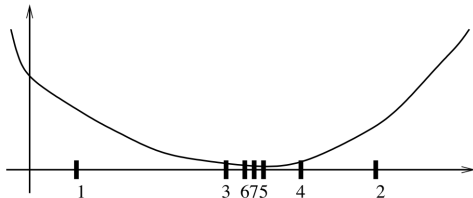
$$(B - A) \min(|f'(A)|, |f'(B)|) \leq \epsilon$$

**Output:**  $x = \frac{A+B}{2}$

solution on the left

# One-dimensional problems

- Key Idea
  - For differentiable  $f$  search for  $\theta$  with  $\nabla f(\theta) = 0$
  - Interval bisection (derivative is monotonic)
- Can be extended to nondifferentiable problems



**Require:**  $a, b$ , Precision  $\epsilon$

Set  $A = a, B = b$

**repeat**

**if**  $f'(\frac{A+B}{2}) > 0$  **then**

$$B = \frac{A+B}{2}$$

**else**

$$A = \frac{A+B}{2}$$

**end if**

**until**



$$(B - A) \min(|f'(A)|, |f'(B)|) \leq \epsilon$$

**Output:**  $x = \frac{A+B}{2}$

solution on the left

# One-dimensional problems

- Key Idea
  - For differentiable  $f$  search for  $\theta$  with  $\nabla f(\theta) = 0$
  - Interval bisection (derivative is monotonic)
- Can be extended to nondifferentiable problems
  - exploit convexity in upper bound and keep 5 points



**Require:**  $a, b$ , Precision  $\epsilon$

Set  $A = a, B = b$

**repeat**

**if**  $f'(\frac{A+B}{2}) > 0$  **then**

$$B = \frac{A+B}{2}$$

**else**

$$A = \frac{A+B}{2}$$

**end if**


**until**

$$(B - A) \min(|f'(A)|, |f'(B)|) \leq \epsilon$$

**Output:**  $x = \frac{A+B}{2}$


solution on the left

# Gradient Descent



- Key Idea

- Gradient points into steepest ascent direction
- Locally, the gradient is a good approximation of the objective function



- GD with Line Search


- Get descent direction, then unconstrained line search
- Turn a multidimensional problem into a one-dimensional problem that we already know how to solve

**given** a starting point  $\theta \in \text{Dom}(f)$

**repeat**

1.  $\Delta \theta := -\nabla f(\theta)$
2. Line search.  $t^* = \arg \min_{t \geq 0} f(\theta + t \cdot \Delta \theta)$
3. Update.  $\theta := \theta + t^* \Delta \theta$

**until** stopping criterion is satisfied.



# Gradient Descent convergence

- Let  $p^*$  be the optimal value,  $\theta^*$  be the minimizer - the point where the minimum is obtained, and  $\theta^{(0)}$  be the starting point
- For strongly convex  $f$  (replace  $\geq$  with  $>$  in the definition of convexity) the residual error  $\rho$ , for the  $k$ -th iteration is:

$$\rho = \underbrace{f(\theta^{(k)}) - p^*}_{\text{residual error}} \leq \underbrace{c^k (f(\theta^{(0)}) - p^*)}_{\text{residual error}}, \quad c < 1$$

$f(\theta^{(k)})$  converges to  $p^*$  as  $k \rightarrow \infty$

- We must have  $f(\theta^{(k)}) - p^* \leq \epsilon$  after at most  $\frac{\log((f(\theta^{(0)}) - p^*)/\epsilon)}{\log(1/c)}$  iterations

*fast converge*

- Linear convergence for strongly convex objective
  - $k \sim \log(\rho^{-1})$  //  $k$  = number of iterations,  $\rho$

0.01  
0.001

- Linear convergence for strongly convex objective
  - i.e. linear when plotting on a log scale - old statistics terminology



# Distributed/Parallel implementation

- Often problems are of the form
  - $f(\boldsymbol{\theta}) = \sum_i L_i(\boldsymbol{\theta}) + g(\boldsymbol{\theta})$
  - where  $i$  iterates over, e.g., each data instance
- Example OLS regression:                      // with regularization
  - $L_i(\mathbf{w}) = (\mathbf{x}_i^T \mathbf{w} - y_i)^2$                        $g(\mathbf{w}) = \lambda \cdot \|\mathbf{w}\|_2^2$
- Gradient can simple be decomposed based on the sum rule
- Easy to parallelize/distribute

# Basic steps


**given** a starting point  $\theta \in \text{Dom}(f)$

**repeat**

1.  $\Delta\theta := -\nabla f(\theta)$
2. Line search.  $t^* = \arg \min_{t>0} f(\theta + t \cdot \Delta\theta)$
3. Update.  $\theta := \theta + t^* \Delta\theta$

**until** stopping criterion is satisfied.

easy parallel computation



evaluating function might be done multiple times: expensive!

- Distribute data over several machines
- Compute partial gradients (on each machine in parallel)
- Aggregate the partial gradients to the final one
- BUT: Line search is expensive
  - for each tested step size: scan through all datapoints

# Scalability analysis

- + Linear time in number of instances
  - + Linear memory consumption in problem size (not data)
  - + Logarithmic time in accuracy
  - + 'Perfect' scalability
- 
- Multiple passes through dataset for each iteration

# A faster algorithm

- Avoid the line search; simply pick update

$$\theta_{t+1} \leftarrow \theta_t - \tau \cdot \nabla f(\theta_t)$$

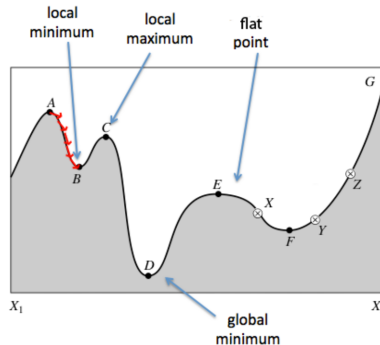
- $\tau$  is often called the **learning rate**

*↑ choose optimal*

- Only a single pass through data per iteration
- Logarithmic iteration bound (as before)
  - if learning rate is chosen "correctly"
- How to pick the learning rate?
  - too small: slow convergence
  - too high: algorithm might oscillate, no convergence
- Interactive tutorial on optimization
  - <http://www.benfrederickson.com/numerical-optimization/>

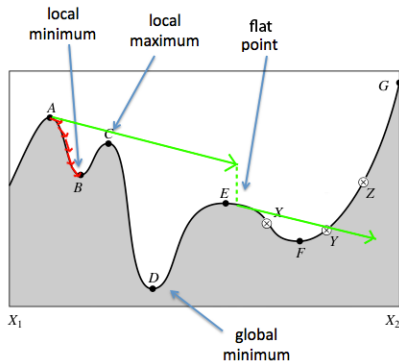
# The value of $\tau$

- A too **small** value for  $\tau$  has two drawbacks
  - We find the minimum more slowly
  - We end up in local minima or saddle/flat points



# The value of $\tau$


- A too **large** value for  $\tau$  has one drawback
  - You may never find a minimum; oscillations usually occur
- We only need 1 step to overshoot



# Learning rate adaptation

- Simple solution: let the learning rate be a decreasing function  $\tau_t$  of the iteration number  $t$ 
  - so called learning rate schedule
  - first iterations cause large changes in the parameters; later do fine-tuning
  - convergence easily guaranteed if  $\lim_{t \rightarrow \infty} \tau_t = 0$
  - example:  $\tau_{t+1} \leftarrow \underline{\alpha} \cdot \tau_t$  for  $0 < \alpha < 1$

# Learning rate adaptation



- Other solutions: Incorporate "history" of previous gradients

- Momentum:

*faster*

- $\mathbf{m}_t \leftarrow \tau \cdot \nabla f(\boldsymbol{\theta}_t) + \gamma \cdot \mathbf{m}_{t-1}$  // often  $\gamma = 0.5$
- $\boldsymbol{\theta}_{t+1} \leftarrow \boldsymbol{\theta}_t - \mathbf{m}_t$
- As long as gradients point to the same direction, the search accelerates

- AdaGrad:

- different learning rate per parameter
- learning rate depends inversely on accumulated "strength" of all previously computed gradients
- large parameter updates ("large" gradients) lead to small learning rates



# Adaptive moment estimation (Adam)



- $\mathbf{m}_t = \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \nabla f(\boldsymbol{\theta}_t)$ 
  - estimate of the first moment (mean) of the gradient
  - Exponentially decaying average of past gradients  $\mathbf{m}_t$  (similar to momentum)
- $\mathbf{v}_t = \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) (\nabla f(\boldsymbol{\theta}_t))^2$ 
  - estimate of the second moment (uncentered variance) of the gradient
  - exponentially decaying average of past squared gradients  $\mathbf{v}_t$
- To avoid bias towards zero (due to 0's initialization) use bias-corrected version instead:
  - $\hat{\mathbf{m}}_t = \frac{\mathbf{m}_t}{1 - \beta_1^t}$        $\hat{\mathbf{v}}_t = \frac{\mathbf{v}_t}{1 - \beta_2^t}$
- Finally, the Adam update rule for parameters  $\boldsymbol{\theta}$ :
  - $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \frac{\tau}{\sqrt{\hat{\mathbf{v}}_t + \epsilon}} \hat{\mathbf{m}}_t$
- Default values:  $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$

# Visualizing gradient descent variants

- AdaGrad and variants

- often have faster convergence
- might help to escape saddlepoints

<http://sebastianruder.com/optimizing-gradient-descent/>



- Gradient descent and similar techniques are called first-order optimization techniques
  - they only exploit information of the gradients (i.e. first order derivative)
- Higher-order techniques use higher-order derivatives
  - e.g. second-order = Hessian matrix
  - Example: Newton Method


# Newton method

- Convex objective function  $f$
- Nonnegative second derivative:  $\nabla^2 f(\theta) \succeq 0$  // Hessian matrix
  - $\nabla^2 f(\theta) \succeq 0$  means that the Hessian is positive semidefinite
- Taylor expansion of  $f$  at point  $\theta_t$

$$f(\theta_t + \delta) = f(\theta_t) + \delta^T \nabla f(\theta_t) + \frac{1}{2} \delta^T \nabla^2 f(\theta_t) \delta + O(\delta^3)$$


$$x^T A x \succeq 0$$

# Newton method



- Convex objective function  $f$
- Nonnegative second derivative:  $\nabla^2 f(\theta) \succeq 0$  // Hessian matrix
  - $\nabla^2 f(\theta) \succeq 0$  means that the Hessian is positive semidefinite

- Taylor expansion of  $f$  at point  $\theta_t$

$$f(\theta_t + \delta) = f(\theta_t) + \delta^T \nabla f(\theta_t) + \frac{1}{2} \delta^T \nabla^2 f(\theta_t) \delta + O(\delta^3) = g(\delta)$$

- Minimize approximation: leads to

$$\theta_{t+1} \leftarrow \theta_t - [\nabla^2 f(\theta_t)]^{-1} \nabla f(\theta_t)$$

$$\begin{aligned} \nabla g(\delta) &= 0 \\ \Leftrightarrow \delta^* &= \text{Minimum} \end{aligned}$$

- Repeat until convergence

# Parallel Newton method

- + Good rate for convergence
- Faster* + Few passes through data needed
- + Parallel aggregation of gradient and Hessian
- + Gradient requires  $O(d)$  data
- Hessian requires  $O(d^2)$  data
- Update step is  $O(d^3)$  & nontrivial to parallelize *To do Inverse  $\nabla^2 f(\theta)^{-1}$*
- Use it only for low dimensional problems!

# Large scale optimization

- Higher-order techniques have nice properties (e.g. convergence) but they are prohibitively expensive for high dimensional problems
- For large scale data / high dimensional problems use first-order techniques *too slow*
  - i.e. variants of gradient descent
- But for real-world large scale data even first-order methods are too costly
- Solution: **Stochastic optimization!**

# Motivation: Stochastic Gradient Descent

- Goal: minimize  $f(\boldsymbol{\theta}) = \sum_{i=1}^n L_i(\boldsymbol{\theta})$  + potential constraints
- For very large data: even a single pass through the data is very costly
- Lots of time required to even compute the very first gradient
- Is it possible to update the parameters more frequently/faster?



# Stochastic Gradient Descent

- Consider the task as empirical risk minimization

$$\frac{1}{n} \left( \sum_{i=1}^n L_i(\boldsymbol{\theta}) \right) = \mathbb{E}_{i \sim \{1, \dots, n\}} [L_i(\boldsymbol{\theta})]$$

- (Exact) expectation can be approximated by smaller sample: = *subset*
- $\mathbb{E}_{i \sim \{1, \dots, n\}} [L_i(\boldsymbol{\theta})] \approx \frac{1}{|S|} \sum_{j \in S} (L_j(\boldsymbol{\theta}))$  // with  $S \subseteq \{1, \dots, n\}$

*subset*

or equivalently:  $\sum_{i=1}^n L_i(\boldsymbol{\theta}) \approx \left[ \frac{n}{|S|} \right] \sum_{j \in S} L_j(\boldsymbol{\theta})$

*↓  
scaling*

# Stochastic Gradient Descent

- Intuition: Instead of using "exact" gradient, compute only a **noisy** (but still **unbiased**) **estimate** based on smaller sample
- Stochastic gradient decent:
  1. randomly pick a (small) subset  $S$  of the points  $\rightarrow$  so called mini-batch
  2. compute gradient based on mini-batch
  3. update:  $\theta_{t+1} \leftarrow \theta_t - \tau \cdot \frac{n}{|S|} \cdot \sum_{j \in S} \nabla L_j(\theta_t)$   $\tau$ : learning rate  
 $\frac{n}{|S|}$ : up-scaling
  4. pick a new subset and repeat with 2
- "Original" SGD uses mini-batches of size 1
  - larger mini-batches lead to more stable gradients (i.e. smaller variance in the estimated gradient)
- In many cases, the data is sampled so that we don't see any data point twice. Then, each full iteration over the complete data set is called one "epoch".

# Example: Perceptron


1. model

2. Loss

- Simple linear binary classifier: ①

$$\delta(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}^T \mathbf{x} + b > 0 \\ -1 & \text{else} \end{cases}$$

- Learning task:  
Given  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$  with  $y_i \in \{-1, 1\}$   
Find  $\min_{\mathbf{w}, b} \sum_i L(y_i, \mathbf{w}^T \mathbf{x}_i + b)$



- $L$  is the loss function, with  $\epsilon > 0$  ②

$$\text{e.g. } L(u, v) = \max(0, \epsilon - u \cdot v) = \begin{cases} \epsilon - uv & \text{if } uv < \epsilon \\ 0 & \text{else } uv \geq \epsilon \end{cases}$$

$u$ : ground truth  
 $v$ : prediction

# Example: Perceptron


- Simple linear binary classifier:

$$\delta(\mathbf{x}) \begin{cases} 1 & \text{if } \mathbf{w}^T \mathbf{x} + b > 0 \\ -1 & \text{else} \end{cases}$$

- Learning task:

Given  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$  with  $y_i \in \{-1, 1\}$

Find  $\min_{\mathbf{w}, b} \sum_i L(y_i, \mathbf{w}^T \mathbf{x}_i + b)$



- $L$  is the loss function, with  $\epsilon > 0$

$u = y_i (\mathbf{w}^T \mathbf{x} + b)$

- e.g.  $L(u, v) = \max(0, \epsilon - u \cdot v) = \begin{cases} \epsilon - uv & \text{if } uv < \epsilon \\ 0 & \text{else} \end{cases}$  ← incorrect prediction  
← correct prediction

$\nabla_{\mathbf{w}} L(y_i, \mathbf{w}^T \mathbf{x}_i + b) = \begin{cases} -y_i \cdot \mathbf{x}_i & \text{if } uv \leq \epsilon \\ 0 & \text{else} \end{cases}$   $\nabla_b L = \begin{cases} -y_i \\ 0 \end{cases}$

# Example: Perceptron

$$\Theta_{i+1} \leftarrow \Theta_i - \tau \cdot \frac{1}{T} \nabla L. \quad \text{if right classify } \nabla L = 0$$

- Let's solve this problem via SGD
- Result:

**initialize**  $w = 0$  and  $b = 0$

**repeat**

**if**  $y_i \cdot (w^T x_i + b) < \epsilon$  **then**

$w \leftarrow w + \tau \cdot n \cdot y_i \cdot x_i$  and  $b \leftarrow b + \tau \cdot n \cdot y_i$

**end if**

**until** all classified correctly

missclassified

$$\tau = \frac{1}{n}$$

- Note: Nothing happens if classified correctly
  - gradient is zero
- Does this remind you of the original learning rules for perceptron?

# Convergence in expectation

- Subject to relatively mild assumptions, stochastic gradient descent converges almost surely to a global minimum when the objective function is convex
  - almost surely to a local minimum for non-convex functions
- The expectation of the residual error decreases with speed

$$\mathbb{E}[\rho] \sim t^{-1}$$

$$// \text{ i.e. } t \sim \mathbb{E}[\rho]^{-1}$$

*residual*  
 $|f(\theta_i) - f(\theta^*)|$   
*iteration.*

- Note: Standard GD has speed  $t \sim \log \rho^{-1}$ 
  - faster convergence speed; but each iteration takes longer

# Optimizing Logistic Regression

- Recall we wanted to solve  $\mathbf{w}^* = \arg \min_{\mathbf{w}} E(\mathbf{w})$

- $E(\mathbf{w}) = -\ln p(\mathbf{y}|\mathbf{w}, \mathbf{X})$  *negative log-likelihood*  
$$= -\sum_{i=1}^N y_i \ln \sigma(\mathbf{w}^T \mathbf{x}_i) + (1 - y_i) \ln(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))$$


- Closed form solution does not exist
- Solution:
  - Compute the gradient  $\nabla E(\mathbf{w})$
  - Find  $\mathbf{w}^*$  using gradient descent
- Is  $E(\mathbf{w})$  convex?
- Can you use SGD?
- How can you choose the learning rate?
- What changes if we add regularization, i.e.  $E_{reg}(\mathbf{w}) = E(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$  ?

# Large-Scale Learning - Distributed Learning

- So far, we (mainly) assumed a single machine
  - SGD achieves speed-up by only operating on a subset of the data
    - Might still be too slow when operating with really large data and large models
  - In practice: We have often multiple machines available
- ⇒ Distributed learning
- Distribute computation across multiple machines
  - Core challenge: distribute work so that communication doesn't kill you



# Distributed Learning: Data vs. Model Parallelism



Use multiple model replicas to process different examples at the same time

- all collaborate to update model state (parameters) in shared parameter server(s)

Many models have lots of inherent parallelism

- local connectivity (as found in CNNs)
- specialized parts of model active only for some examples (see, e.g., Matrix Factorization)

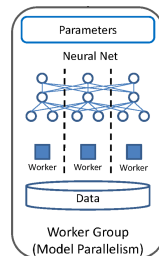


figure based on <https://svn.apache.org/repos/infra/websites/production/singa/content/v0.1.0/architecture.html>

# Parameter Server

- General goal: Keep time to send/receive parameters over network small, compared to the actual time used for computation




figure from *Large Scale Distributed Systems for Training Neural Networks*, Jeff Dean

# Distributed Learning in Practice

- Distributed optimization/learning is essential when operating with very large data (and large models)
  - Default for training ML models in today's production systems
- Many modern ML frameworks (e.g. Tensorflow, PyTorch, MXNet, ...) provide support for distributed learning
- Many further aspects/challenges
  - Desired synchronization
  - Fault tolerance, recovery
  - Automatic placement (of data/model) to reduce communication

# Summary

- General task: Find solution  $\theta^*$  minimizing function  $f$
- Convex sets & functions
  - Global vs. local minimum
  - Verifying convexity: Definition, special results (first-order convexity, 2nd derivative), convexity-preserving operations
- Gradient descent:  $\theta := \theta - t \nabla f(\theta)$ 
  - How to choose  $t$ ? Line search, fixed
  - Learning rate: Fix  $t = \tau$ ; or use an adaptive learning rate (momentum, AdaGrad, Adam)
  - Stochastic gradient descent (SGD): Only use part of data (mini-batches) at each step
- Distributed Learning: exploit multiple machines
  - data parallelism, model parallelism

# Reading material

## Reading material

- Boyd - Convex Optimization: chapters 2.1-2.3, 3.1, 3.2, 4.1, 4.2, 9
  - free PDF version online
- Sebastian Ruder - An overview of gradient descent optimization algorithms
  - <https://arxiv.org/abs/1609.04747>