


Machine Learning for Graphs and Sequential Data Exercise Sheet 6

Graphs: Clustering

Problem 1: Given the graph below, find the following partitionings of the graph for $k = 2$:

- The partitioning giving the global minimum cut
- A partitioning approximately minimizing the ratio cut
- A partitioning approximately minimizing the normalized cut



```
import itertools

import numpy as np
from scipy.linalg import eigh
from scipy.sparse import coo_matrix

def laplacian(A):
    D = A.sum(axis=0)
    return np.identity(A.shape[0]) * D - A

def format_partitioning(f):
    c1 = (f > 0.0).nonzero()[0]
    c2 = (f < 0.0).nonzero()[0]
    return f"{{{',_'.join(map(str,_c1))}}}_{{{',_'.join(map(str,_c2))}}}"

edges = np.array([
    (0, 1, 2), (0, 3, 2), (1, 2, 4), (1, 3, 2), (2, 3, 1),
    (2, 4, 3), (2, 5, 4), (2, 6, 4), (4, 5, 2), (5, 6, 1)])
A = coo_matrix((edges[2], (edges[0], edges[1])), shape=(7, 7))
A = A.toarray()
A = A + A.T
L = laplacian(A)

# Global minimum cut
```

```

fs = [((f := np.array(f_)) @ L @ f / 4, f)
       for f_ in itertools.product([-1, 1], repeat=A.shape[0])]
if len(set(f_)) > 1]
min_cost, min_cut = min(fs, key=lambda f: f[0])
print(f"Global_minimum_cut_is_{format_partitioning(min_cut)}_at_cost_{min_cost}")

# Approximate ratio cut
lambda_, v = eigh(L, eigvals=(1, 1))
print(f"Approximate_ratio_cut_is_{format_partitioning(v.squeeze())}")

# Approximate normalized cut
D_sqrt_inv = np.diag(1 / np.sqrt(A.sum(axis=1)))
L_normalized = D_sqrt_inv @ L @ D_sqrt_inv
lambda_, v = eigh(L_normalized, eigvals=(1, 1))
print(f"Approximate_normalized_cut_is_{format_partitioning(v.squeeze())}")

```

The above solution outputs

Global minimum cut is {1, 2, 3, 4, 5, 6} {0} at cost 4.0
 Approximate ratio cut is {0, 1, 3} {2, 4, 5, 6}
 Approximate normalized cut is {2, 4, 5, 6} {0, 1, 3}

Problem 2: Consider minimizing the ratio cut on a graph with two clusters C_1 and C_2 and N nodes in total. The indicator vector

$$f_{C_1, i} = \begin{cases} +\sqrt{\frac{|C_1|}{|C_1|}} & \text{if } v_i \in C_1 \\ -\sqrt{\frac{|C_1|}{|C_1|}} & \text{otherwise} \end{cases}$$

is defined as in the lecture. Prove the following three properties about \mathbf{f}_{C_1} .

- $1^T \mathbf{f}_{C_1} = \sum_i f_{C_1, i} = 0$
- $\mathbf{f}_{C_1}^T \mathbf{f}_{C_1} = \|\mathbf{f}_{C_1}\|_2^2 = |V|$
- $\mathbf{f}_{C_1}^T L \mathbf{f}_{C_1} = |V| \left[\frac{\text{cut}(C_1, C_2)}{|C_1|} + \frac{\text{cut}(C_1, C_2)}{|C_1|} \right]$

Let $M = |C_1|$.

a)

$$\begin{aligned}
1^T \mathbf{f}_{C_1} &= \sum_i f_{C_1,i} \\
&= M \sqrt{\frac{|C_1|}{|C_1|}} - (N-M) \sqrt{\frac{|C_1|}{|C_1|}} \\
&= M \sqrt{\frac{N-M}{M}} - (N-M) \sqrt{\frac{M}{N-M}} \\
&= \sqrt{M(N-M)} - \sqrt{(N-M)M} = 0
\end{aligned}$$

b)

$$\begin{aligned}
\|\mathbf{f}_{C_1}\|_2^2 &= \mathbf{f}_{C_1}^T \mathbf{f}_{C_1} = \sum_i f_{C_1,i}^2 \\
&= \sum_i \begin{cases} \frac{|C_1|}{|C_1|} & \text{if } v_i \in C_1 \\ \frac{|C_1|}{|C_1|} & \text{otherwise} \end{cases} \\
&= M \cdot \frac{N-M}{M} + (N-M) \cdot \frac{M}{N-M} = N = |V|
\end{aligned}$$

c)

$$\begin{aligned}
\mathbf{f}_{C_1}^T L \mathbf{f}_{C_1} &= \frac{1}{2} \sum_{(u,v) \in E} W_{uv} (f_{C_1,u} - f_{C_1,v})^2 && \text{(Observation 1)} \\
&= \frac{1}{2} \left(\sum_{\substack{(u,v) \in E \\ u,v \in C_1}} W_{uv} (f_{C_1,u} - f_{C_1,v})^2 + \sum_{\substack{(u,v) \in E \\ u,v \in C_2}} W_{uv} (f_{C_1,u} - f_{C_1,v})^2 \right. \\
&\quad \left. + 2 \sum_{\substack{(u,v) \in E \\ u \in C_1, v \in C_2}} W_{uv} (f_{C_1,u} - f_{C_1,v})^2 \right)
\end{aligned}$$


The first two terms are 0 because $f_{C_1,u} = f_{C_1,v}$ if u and v are in the same cluster.

$$\begin{aligned}
 &= \sum_{\substack{(u,v) \in E \\ u \in C_1, v \in C_2}} W_{uv} (f_{C_1,u} - f_{C_1,v})^2 \\
 &= \sum_{\substack{(u,v) \in E \\ u \in C_1, v \in C_2}} W_{uv} (f_{C_1,u}^2 - 2f_{C_1,u}f_{C_1,v} + f_{C_1,v}^2) \\
 &= \sum_{\substack{(u,v) \in E \\ u \in C_1, v \in C_2}} W_{uv} \left(\frac{N-M}{M} + 2\sqrt{\frac{(N-M)M}{M(N-M)}} + \frac{M}{N-M} \right) \\
 &= \sum_{\substack{(u,v) \in E \\ u \in C_1, v \in C_2}} W_{uv} \left(\frac{N-M}{M} + 2 + \frac{M}{N-M} \right) \\
 &= \left(\frac{N-M}{M} + 2 + \frac{M}{N-M} \right) \sum_{\substack{(u,v) \in E \\ u \in C_1, v \in C_2}} W_{uv}
 \end{aligned}$$

The sum of all edge weights of edges with one end in C_1 and the other in C_2 is exactly the size of the cut

$$\begin{aligned}
 &= \left(\frac{N-M}{M} + 2 + \frac{M}{N-M} \right) \text{cut}(C_1, C_2) \\
 &= \left(\frac{N-M}{M} + \frac{M}{M} + \frac{N-M}{N-M} + \frac{M}{N-M} \right) \text{cut}(C_1, C_2) \\
 &= \left(\frac{N}{M} + \frac{N}{N-M} \right) \text{cut}(C_1, C_2) \\
 &= |V| \left[\frac{\text{cut}(C_1, C_2)}{|C_1|} + \frac{\text{cut}(C_1, C_2)}{|C_2|} \right]
 \end{aligned}$$

Problem 3: Answer the following questions regarding the graph below. Formulate a conjecture first and then verify it computationally in a notebook.



- a) How does the first eigenvector change when increasing the weight between node v_6 and v_9 ?
- b) How does the spectral embedding change?
- c) How does this change affect the final clustering?

The first eigenvector will not change because it is always the all-ones vectors scaled to have length 1. The embeddings of node 6 and 9 will move closer together because their connection becomes stronger, other nodes will adapt accordingly. Depending on the increase, node 9 might change into the same cluster as node 6.

We verify our hypotheses with the following program which produces ??.

```
import matplotlib as mpl
mpl.use("Agg")
import matplotlib.pyplot as pp

import numpy as np
from scipy.linalg import eigh
from scipy.sparse import coo_matrix
from sklearn.cluster import KMeans
import seaborn as sns

sns.set()

def plot_cluster(ebds, epsilon=0.05):
    # Rotate embeddings such that node 4 is always embedded straight down
    four = ebds[3]
    alpha = np.arccos(-four[1] / np.linalg.norm(four))
    c, s = np.cos(alpha), np.sin(alpha)
    R = np.array([[c, -s], [s, c]])
    ebds = ebds @ R

    clusters = KMeans(n_clusters=3).fit_predict(ebds)

    fig, ax = pp.subplots(1, 1, figsize=(6, 6))
    xtp, ytp = ebds.ptp(axis=0) * 0.1
    ax.set_xlim((ebds[:, 0].min() - xtp, ebds[:, 0].max() + xtp))
    ax.set_ylim((ebds[:, 1].min() - ytp, ebds[:, 1].max() + ytp))

    # Disturb points to show nodes that get mapped to the same coordinates
    points = ebds + np.random.rand(*ebds.shape) * epsilon

    colors = ["firebrick", "seagreen", "dodgerblue"]
    bbox_props = dict(boxstyle="circle", alpha=0.5, ec="b", lw=2)
    for i, xyc in enumerate(zip(points, clusters)):
        xy, c = xyc
        bbox_props["fc"] = colors[c]
```

```
        pp.text(xy[0], xy[1], str(i + 1), bbox=bbox_props,
                horizontalalignment="center", verticalalignment="center")

    return fig, ax


def laplacian(A):
    D = A.sum(axis=0)
    return np.identity(A.shape[0]) * D - A

edges = np.array([
    (1, 2, 1), (1, 3, 1), (2, 3, 1), (2, 5, 1), (4, 5, 1),
    (4, 6, 1), (5, 6, 2), (6, 9, 1), (7, 8, 3), (7, 9, 1), (8, 9, 1)]).T
edges[:2] -= 1
A = coo_matrix((edges[2], (edges[0], edges[1])), shape=(9, 9)).toarray()
A = A + A.T

_, eig = eigh(laplacian(A), eigvals=(1, 2))
fig, _ = plot_cluster(eig)
fig.savefig("problem_3_pre.png")

A[5, 8] = 4
A[8, 5] = 4
_, eig = eigh(laplacian(A), eigvals=(1, 2))

fig, _ = plot_cluster(eig)
fig.savefig("problem_3_post.png")
```



(a) Original weights

(b) After increasing weight between nodes 6 and 9

Figure 1: Spectral embeddings in problem 3