

Machine Learning for Graphs and Sequential Data

Robustness of Machine Learning – Exact Certification

Lecturer: Prof. Dr. Stephan Günnemann

cs.cit.tum.de/daml

Summer Term 2023

Data Analytics and
Machine Learning



Roadmap

1. Introduction
2. Construction of adversarial examples
3. Improving robustness
- 4. Certifiable robustness**
 - **Exact certification**
 - Convex relaxations
 - Lipschitz-continuity
 - Randomized smoothing

对抗性训练提高了鲁棒性，但我们如何才能确保用户能够真正依赖结果？
- 仍然可能存在某些情况下，模型的行为是不受欢迎的

正如所讨论的，对对抗性例子的检测似乎并不奏效

更好的方法：稳健性认证。

- 思路：尝试证明分类器的预测在一个半径范围内没有变化（用某个规范来衡量）。
- 如果证明成功，我们就知道在这个半径内不可能有对抗性的例子。我们得到了一个保证！
- 如果证明不成功，预测就会改变。因此，样本可能是一个对抗性的例子（或者有可能对抗性地改变它）。
- 在一个非常保守的方法中，我们现在可以拒绝我们模型的预测和/或咨询专家进行人工检查。

Motivation

- Adversarial Training improves robustness, but how can we make sure that the user can really rely on the results?
 - There could still exist some cases where the model behaves in an undesired way
- As discussed, detection of adversarial examples does not seem to work
- Better approach: Robustness certification.
 - Idea: try to prove that the classifier's prediction does not change within a radius (measured by some norm)
 - If the proof is successful, we know there cannot be an adversarial example within that radius. We get a guarantee!
 - If the proof is not successful, the prediction could change. Therefore the sample might be an adversarial example (or it might be possible to adversarially change it).
 - In a very conservative approach, we could now refuse our model's prediction and/or consult an expert for manual inspection.

Exact Verification

目标：开发一种能回答问题的算法：

"分类器 f ！在样本 \mathbf{x} 周围是否无对抗性（在某个规范测量的 ϵ 球内）？"

当且仅当输入样本周围的 ϵ 球内没有对抗性例子时，该算法应返回 YES（即 NO，如果有一个对抗性例子）。

精确验证方法通常是具有 ReLU 激活函数的神经网络设计的。ReLU 网络在深度学习中非常普遍，非常适合于组合式精确验证方法。

Goal: Develop an algorithm that answers the question:

"Is the classifier f_θ around the sample \mathbf{x} adversarial-free
(within an ϵ -ball measured by some norm)?"

The algorithm should return **YES** if and only if there are no adversarial example
within an ϵ ball around the input sample (i.e. **NO** iff there is an adv. example).

Exact verification methods are typically designed for neural networks with **ReLU activation function**. ReLU networks are very prevalent in deep learning and are well-suited for **combinatorial** exact verification methods.

Exact Verification

- We view the neural network as a **sequence of functions** (i.e. the layers).
- Each **layer** is defined as $f_i(x) = \sigma(\mathbf{W}_i x + b_i)$, where \mathbf{W}_i and b_i are the weight matrix and the bias of layer i , respectively.
- The **ReLU activation** function is defined as $\sigma(x) = \max(0, x)$ and is applied entry-wise to the input.
- The **overall network** is a function $F: \mathbb{R}^d \rightarrow \mathbb{R}^{|y|}$ given by:
$$F(x) = \mathbf{W}_L f_{L-1} \circ f_{L-2} \circ \dots \circ f_1(x) + \mathbf{b}_L$$
- The output of F are the **logits** which are subsequently fed into the softmax function to obtain a categorical distribution.
 - We can omit the softmax for certification since the operation is order-preserving (i.e., the ‘winning’ class does not change).

Exact Certification: Complexity

Exact certification is a very powerful method for a defending system: we know exactly when a sample could be an adversarial example and can potentially even use this knowledge to get the worst-case perturbation for adversarial training.

Unfortunately, [Katz et al. 2017] report the following result:

Theorem: Exact certification of neural networks with ReLU activation function and L_∞ -bounded perturbations is **NP-complete**.

Nevertheless, solvers for NP-complete problems have made significant progress, so certifying small to medium-size neural networks is sometimes possible.

准确认证对于防卫系统来说是一个非常强大的方法：我们确切地知道一个样本什么时候可能是一个对抗性的例子，甚至有可能利用这一知识来获得对抗性训练的最坏情况下的扰动。

不幸的是，[Katz等人，2017]报告了以下结果：

定理：具有ReLU激活函数和 L^* 有界扰动的神经网络的精确认证是NP-complete。

尽管如此，NP-complete问题的求解器已经取得了重大进展，所以认证中小型神经网络有时是可能的。

Mixed Integer Linear Programming

- One approach for exact certification of ReLU networks is to use mixed integer linear programming (**MILP**).
- Recall **linear programs (LPs)**:

minimize
subject to

$$\mathbf{c}^T \mathbf{x}$$

$$\mathbf{Ax} \leq \mathbf{b}$$
$$\mathbf{x} \geq 0$$

Linear objective

Linear subject

- Integer linear programs: we have the additional constraints $\mathbf{x}_i \in \mathbb{Z}$, i.e. the variables are integer-valued
- Mixed integer linear programs (**MILP**): *some* variables constrained to be integers, others not.

Mixed Integer Linear Programming: Complexity

- One approach for exact certification of ReLU networks is to use mixed integer linear programming (**MILP**).

- Recall **linear programs** (LPs):

$$\begin{array}{ll} \text{minimize} & \mathbf{c}^T \mathbf{x} \\ \text{subject to} & \mathbf{Ax} \leq \mathbf{b} \\ & \mathbf{x} \geq 0 \end{array}$$

Can be solved **efficiently**
(in polynomial time)

- Integer linear programs: we have the additional constraints $\mathbf{x}_i \in \mathbb{Z}$, i.e. the variables are integer-valued
- Mixed integer linear programs (**MILP**): *some* variables constrained to be integers, others not.

NP-complete

Expressing Exact Certification as MILP

- Suppose our classifier predicts class c^* for \mathbf{x} , i.e. $c^* = \arg \max_c F(\mathbf{x})_c$
- We call $m_t = F(\mathbf{x})_{c^*} - F(\mathbf{x})_t$ the classification margin of classes c^* and t .
margin GT class t
- Worst-case margin:

$$m_t^* = \min_{\tilde{\mathbf{x}}} F(\tilde{\mathbf{x}})_{c^*} - F(\tilde{\mathbf{x}})_t$$


subject to $\|\tilde{\mathbf{x}} - \mathbf{x}\|_p \leq \epsilon$

- $m_t^* > 0$: the classifier's prediction cannot be changed from class c^* to t
- If for **all classes** $t \neq c^*$ we have $m_t^* > 0 \rightarrow$ we can certify robustness
- If for **any class** $t \neq c^*$ we have $m_t^* < 0 \rightarrow$ there exists an adversarial example $\tilde{\mathbf{x}} \in \mathcal{P}_{\epsilon,p}(\mathbf{x})$

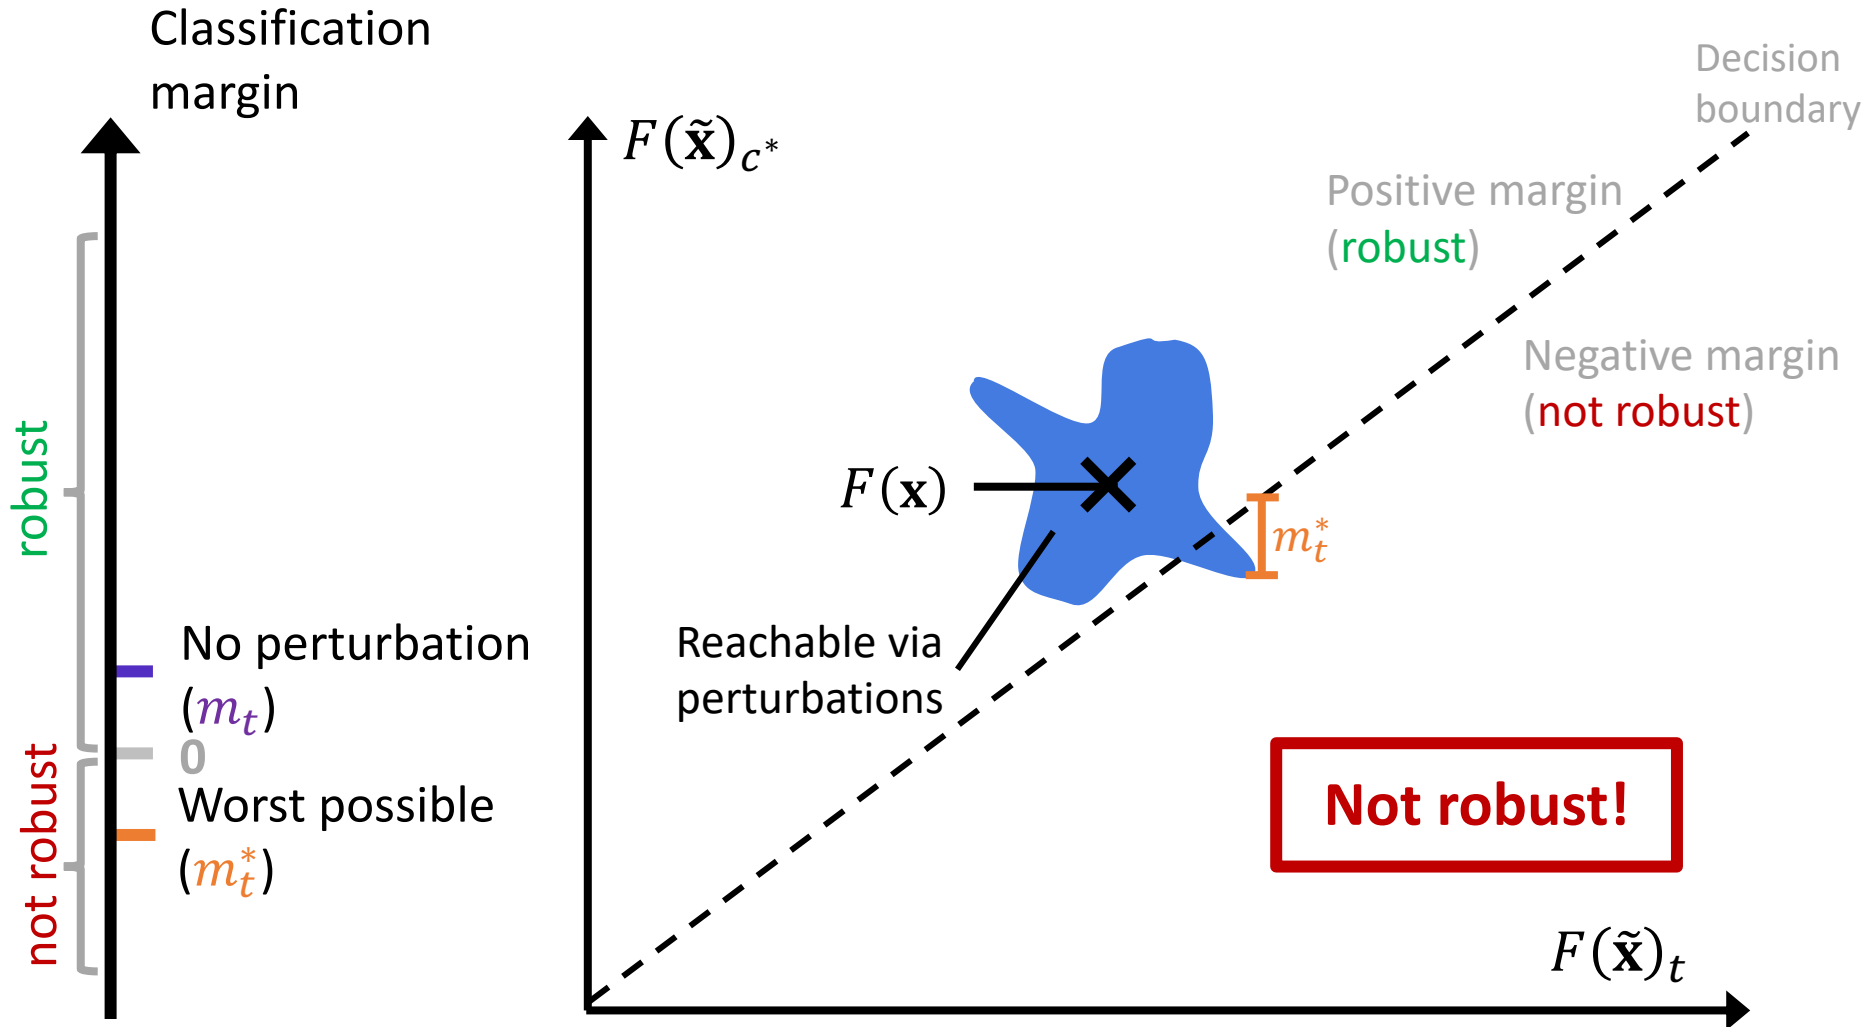
*证分类
对 $t \rightarrow$ 飞机*

$F(\tilde{\mathbf{x}})_{c^} < F(\tilde{\mathbf{x}})_t$*

Exact Robustness Certification: Illustration



Exact Robustness Certification: Illustration



Exact Certification: Optimization Problem

- We can write the optimization problem:

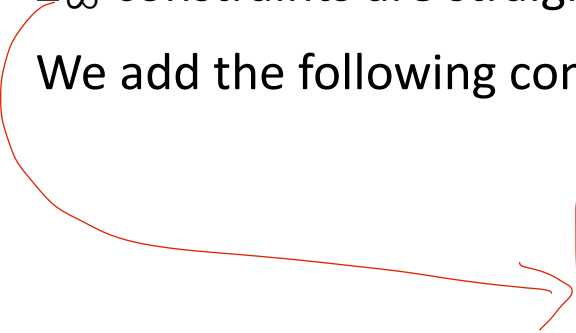
$$\begin{aligned}
 m_t^* &= \min_{\tilde{\mathbf{x}}, \mathbf{y}, \hat{\mathbf{x}}^{(l)}} [\hat{\mathbf{x}}^{(L)}]_c^* - [\hat{\mathbf{x}}^{(L)}]_t && \leftarrow \text{Linear.} \\
 \text{subject to } & \|\tilde{\mathbf{x}} - \mathbf{x}\|_p \leq \epsilon && \leftarrow \text{Linear} \\
 & \mathbf{y}^{(0)} = \tilde{\mathbf{x}} \\
 & \hat{\mathbf{x}}^{(l)} = \mathbf{W}_l \mathbf{y}^{(l-1)} + \mathbf{b}_l && \forall l = 1 \dots L \\
 & \mathbf{y}^{(l)} = \text{ReLU}(\hat{\mathbf{x}}^{(l)}) && \forall l = 1 \dots L - 1
 \end{aligned}$$

non linear →

- Here, $\hat{\mathbf{x}}^{(l)}$ denotes the pre-ReLU activation at layer l .
- To express this as a MILP, we need to encode
 - The L_p constraints on the adversarial perturbation
 - The nonlinear ReLU constraints, which is where most of the difficulty comes from

Expressing Exact Certification as MILP: L_p Constraints

- L_∞ constraints are straightforward to include in linear programs
- We add the following constraints to the optimization:


$$\begin{aligned} \mathbf{x}_i - \tilde{\mathbf{x}}_i &\leq \epsilon \quad \forall i \\ \tilde{\mathbf{x}}_i - \mathbf{x}_i &\leq \epsilon \quad \forall i \end{aligned}$$




- L_1 constraints are also straightforward to encode
- L_2 constraints can be captured using mixed integer quadratic programming

Expressing Exact Certification as MILP: ReLU (1)

- The ReLU activation function is where most of the difficulty comes from.
- We want to encode $\mathbf{y} = \text{ReLU}(\mathbf{x})$
- Naively, we can encode the ReLU activation by introducing a binary variable (vector of variables) \mathbf{a} : $\in \{0, 1\}$

$$\mathbf{y}_i \leq \mathbf{a}_i \cdot \mathbf{x}_i \text{ and } \mathbf{y}_i \geq 0 \text{ and } \mathbf{y}_i \geq \mathbf{x}_i = \text{ReLU}$$

- However, now we have a constraint with a product of two variables, hence, is not linear and not convex.



Expressing Exact Certification as MILP: ReLU (2)

- Suppose we have lower and upper bounds $[l, u]$ on the input x to the ReLU activation.
- Then, we can encode the ReLU activation using linear and integer constraints [Tjeng et al. 2019]):


$$y_i = \text{ReLU}(x_i) \Leftrightarrow \begin{aligned} & (y_i \leq x_i - l_i(1 - a_i)) \\ & \wedge (y_i \leq a_i \cdot u_i) \\ & \wedge (y_i \geq x_i) \\ & \wedge (y_i \geq 0) \\ & \wedge (a_i \in \{0, 1\}) \end{aligned} \quad \forall i$$

Handwritten notes for $a=0$:

$$\begin{aligned} y &\leq x - l \\ y &\leq 0 \\ \hline y &\geq 0 \end{aligned}$$

Handwritten notes for $a=1$:

$$\begin{aligned} y &\leq x \\ y &\leq 1 \cdot u \end{aligned}$$



Overall MILP

- Note that the overall MILP optimizes over multiple variables
 - Efficient solvers will remove some of them due to the equality constraints

$$m_t^* = \min_{\tilde{\mathbf{x}}, \mathbf{y}^{(l)}, \hat{\mathbf{x}}^{(l)}, \mathbf{a}^{(l)}} [\hat{\mathbf{x}}^{(L)}]_{c^*} - [\hat{\mathbf{x}}^{(L)}]_t$$

subject to $\mathbf{x}_i - \tilde{\mathbf{x}}_i \leq \epsilon \quad \forall i$
 $\tilde{\mathbf{x}}_i - \mathbf{x}_i \leq \epsilon \quad \forall i$ } small perturbations

$$\mathbf{y}^{(0)} = \tilde{\mathbf{x}}$$

$$\hat{\mathbf{x}}^{(l)} = \mathbf{W}_l \mathbf{y}^{(l-1)} + \mathbf{b}_l \quad \forall l = 1 \dots L$$

ReLU

$$\left[\begin{array}{l} \mathbf{y}_i^{(l)} \leq \hat{\mathbf{x}}_i^{(l)} - \mathbf{l}_i^{(l)} (1 - \mathbf{a}_i^{(l)}) \\ \mathbf{y}_i^{(l)} \geq \hat{\mathbf{x}}_i^{(l)} \\ \mathbf{y}_i^{(l)} \leq \mathbf{u}_i^{(l)} \cdot \mathbf{a}_i^{(l)} \\ \mathbf{y}_i^{(l)} \geq 0 \\ \mathbf{a}_i^{(l)} \in \{0, 1\} \end{array} \right. \quad \forall l = 1 \dots L - 1 \quad \forall i$$

- Remark: We can also handle all classes $t \neq c^*$ at the same time in a single MILP

On the Lower and Upper Bounds

- The MILP formulation relies on being able to compute lower and upper bounds $[l^{(l)}, u^{(l)}]$ on the input $\hat{\mathbf{x}}^{(l)}$ to the ReLU activation (for every layer l)
- One simple way to get (loose) lower and upper bounds is interval arithmetic

$$\left\{ \begin{array}{l} u^{(l)} = [W_l]_+ u^{(l-1)} - [W_l]_- l^{(l-1)} + b_l \\ l^{(l)} = [W_l]_+ l^{(l-1)} - [W_l]_- u^{(l-1)} + b_l \end{array} \right.$$

positive upper negative lower.
positive lower negative upper

L_p constraints

$$u_i^{(0)} = x_i + \epsilon$$

$$l_i^{(0)} = x_i - \epsilon$$

- Here, $[W]_+ = \max \{W, 0\}$, $[W]_- = \max \{-W, 0\}$

Stable and Unstable Units

- Units for which $\mathbf{u}_i^{(l)} \geq \mathbf{l}_i^{(l)} \geq 0$ are called stably active $y_i = x$
 - Units for which $0 \geq \mathbf{u}_i^{(l)} \geq \mathbf{l}_i^{(l)}$ are called stably inactive $y_i = 0$
 - Units for which $\mathbf{u}_i^{(l)} \geq 0 \geq \mathbf{l}_i^{(l)}$ are called unstable
- } Can be removed from the optimization
- While the tightness of the upper and lower bounds has no influence on the correctness of the result, tighter bounds lead to more stable units and greatly speed up the optimization.

虽然上下限的严密性对结果的正确性没有影响，但严密的界限会导致更稳定的单元，并大大加快优化的速度。

Summary

- Exact verification is possible for ReLU-Networks
 - However, expensive for large neural networks
- Can we find more efficient certificates?
 - Unfortunately not if we focus on exact certificates ("if and only if") due to the NP-hardness (assuming $P \neq NP$)
- However, we can change the allowed answers to our question

“Is the classifier f_θ around the sample \mathbf{x} adversarial-free
(within an ϵ -ball measured by some norm)?”
- If the algorithm returns **YES**, there are no adversarial examples within an ϵ ball around the input sample; if the algorithm returns **POTENTIALLY NOT** there might be adversarial examples or it is adversarial-free *Idk*
- This is a conservative (careful) answer, i.e. in cases where the algorithms says “yes” we can rely on the prediction → i.e. we still have a guarantee
- We discuss such principles in the next sections!

Recommended Reading

- Lecture 12: Certified defenses I: Exact certification of Jerry Li's course on Robustness in Machine Learning (CSE 599-M), <https://jerryzli.github.io/robust-ml-fall19.html>

References – Exact Verification

- Katz, Guy, et al. "Reluplex: An efficient SMT solver for verifying deep neural networks." *International Conference on Computer Aided Verification*. Springer, Cham, 2017.
- Tjeng, Vincent, et al. "Evaluating Robustness of Neural Networks with Mixed Integer Programming." *International Conference on Learning Representations*, 2019