

Machine Learning for Graphs and Sequential Data Exercise Sheet 07

Robustness of Machine Learning Models

Exercises marked with a (*) will be discussed in the in-person exercise session.

Problem 1: (*) Suppose we have a trained binary logistic regression classifier with weight vector $\mathbf{w} \in \mathbb{R}^d$ and bias $b \in \mathbb{R}$. Given a sample $\mathbf{x} \in \mathbb{R}^d$ we want to construct an adversarial example via gradient descent on the binary cross entropy loss:

$$\mathcal{L}(\mathbf{x}, y) = -y \log(\sigma(z)) - (1 - y) \log(1 - \sigma(z)),$$

where $\sigma(z) = \frac{1}{1+e^{-z}}$ is the logistic sigmoid function, $z = \mathbf{w}^T \mathbf{x} + b$, and $y \in \{0, 1\}$ is the class label of the sample at hand.

- a) Derive the gradient $\nabla_{\mathbf{x}} \mathcal{L}(\mathbf{x}, y)$. How do you interpret the result?

Hint: You may use the relation $1 - \sigma(z) = \sigma(-z)$.

$$\begin{aligned} \nabla_{\mathbf{x}} \mathcal{L}(\mathbf{x}, y) &= \frac{-y}{\sigma(z)} \frac{\partial \sigma(z)}{\partial z} \nabla_{\mathbf{x}} z - \frac{1-y}{\sigma(-z)} \frac{\partial \sigma(-z)}{\partial z} \nabla_{\mathbf{x}} z \\ &= \frac{-y}{\sigma(z)} \sigma(z) \sigma(-z) \mathbf{w} + \frac{1-y}{\sigma(-z)} \sigma(-z) \sigma(z) \mathbf{w} \\ &= -y \sigma(-z) \mathbf{w} + (1-y) \sigma(z) \mathbf{w} \end{aligned}$$

The gradient is orthogonal to the decision boundary and points in the direction of the wrong class, depending on y .

- b) Provide a closed-form expression for the worst-case perturbed instance $\tilde{\mathbf{x}}^*$ (measured by the loss \mathcal{L}) for the perturbation set $\mathcal{P}(\mathbf{x}) = \{\tilde{\mathbf{x}} : \|\tilde{\mathbf{x}} - \mathbf{x}\|_2 \leq \epsilon\}$, i.e.

$$\tilde{\mathbf{x}}^* = \arg \max_{\|\tilde{\mathbf{x}} - \mathbf{x}\|_2 \leq \epsilon} \mathcal{L}(\tilde{\mathbf{x}}, y)$$

Since the loss is convex w.r.t. the data, taking a gradient step of magnitude ϵ towards the wrong class will result in the maximum increase in loss:

$$\begin{aligned} \tilde{\mathbf{x}}^* &= \mathbf{x} - \epsilon \frac{\mathbf{w}}{\|\mathbf{w}\|_2} \text{ if } y = 1 \\ \tilde{\mathbf{x}}^* &= \mathbf{x} + \epsilon \frac{\mathbf{w}}{\|\mathbf{w}\|_2} \text{ if } y = 0 \end{aligned}$$

- c) What is the smallest value of ϵ for which the sample \mathbf{x} is misclassified (assuming it was correctly classified before)?

For the sample to change classification we need to have $\sigma(z) = 0.5 \Leftrightarrow \mathbf{w}^T \tilde{\mathbf{x}} + b = 0$. Plugging in the perturbation we get for $y = 1$:

$$\begin{aligned}\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \epsilon \frac{\mathbf{w}}{\|\mathbf{w}\|_2} + b &= 0 \\ \mathbf{w}^T \mathbf{x} - \epsilon \|\mathbf{w}\|_2 + b &= 0 \\ \frac{1}{\|\mathbf{w}\|_2} (\mathbf{w}^T \mathbf{x} + b) &= \epsilon\end{aligned}$$

Thus, for a misclassification we need $\epsilon > \frac{1}{\|\mathbf{w}\|_2} (\mathbf{w}^T \mathbf{x} + b)$.

Analogously for $y = 0$ we obtain $\epsilon > \frac{1}{\|\mathbf{w}\|_2} (-\mathbf{w}^T \mathbf{x} - b)$

- d) We would now like to perform adversarial training. Provide a closed-form expression of the worst-case loss

$$\hat{\mathcal{L}}(\mathbf{x}, y) = \max_{\|\tilde{\mathbf{x}} - \mathbf{x}\|_2 \leq \epsilon} \mathcal{L}(\tilde{\mathbf{x}}, y)$$

as a function of \mathbf{x} and \mathbf{w} . How do you interpret the results?

$$\begin{aligned}\hat{\mathcal{L}}(\mathbf{x}, y) &= \max_{\|\tilde{\mathbf{x}} - \mathbf{x}\|_2 \leq \epsilon} \mathcal{L}(\tilde{\mathbf{x}}, y) \\ &= \mathcal{L}(\tilde{\mathbf{x}}^*, y) \\ &= -y \log(\sigma(\mathbf{w}^T \mathbf{x} - \epsilon \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|_2} + b)) - (1 - y) \log(\sigma(-\mathbf{w}^T \mathbf{x} - \epsilon \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|_2} - b)) \\ &= -y \log(\sigma(\mathbf{w}^T \mathbf{x} - \epsilon \|\mathbf{w}\|_2 + b)) - (1 - y) \log(\sigma(-\mathbf{w}^T \mathbf{x} - \epsilon \|\mathbf{w}\|_2 - b))\end{aligned}$$

Consider the case $y = 1$ ($y = 0$ follows symmetrically). The input to the sigmoid function is shifted to the left (i.e. negative direction) by $\epsilon \|\mathbf{w}\|_2$, reducing the predicted probability of the sample \mathbf{x} belonging to class 1. Thus, only if $\mathbf{w}^T \mathbf{x} + b \geq \epsilon \|\mathbf{w}\|_2$ the sample will be classified as belonging to class 1. We can interpret this as trying to enforce that each sample has at least a distance of $\epsilon \|\mathbf{w}\|_2$ to the decision boundary. Moreover, this margin is proportional to the norm of the weight vector, so simply increasing the norm of \mathbf{w} does not lead to the desired outcome, since we can move $\epsilon \|\mathbf{w}\|_2$ units towards the decision boundary for a unit norm change on the sample \mathbf{x} . Note that, in contrast to support vector machines (SVMs), even when the samples have a margin of at least $\epsilon \|\mathbf{w}\|_2$ to the decision boundary, we have non-zero loss and continue training.

Problem 2: (*) In the lecture on exact certification of neural network robustness we have considered $K - 1$ optimization problems (one for each incorrect class) of the form (c.f. slide 42):

$$m_t^* = \min_{\tilde{\mathbf{x}}, \mathbf{y}^{(t)}, \hat{\mathbf{x}}^{(t)}, \mathbf{a}^{(t)}} [\hat{\mathbf{x}}^{(L)}]_{c^*} - [\hat{\mathbf{x}}^{(L)}]_t \quad \text{subject to MILP constraints.}$$

That is, for each class $t \neq c^*$, we optimize for the **worst-case margin** m_t^* , and conclude that the classifier is robust if and only if

$$\min_{t \neq c^*} m_t^* \geq 0.$$

However, we can equivalently solve the following single optimization problem:

$$m^* = \min_{\hat{\mathbf{x}}, \mathbf{y}^{(l)}, \hat{\mathbf{x}}^{(l)}, \mathbf{a}^{(l)}} \left([\hat{\mathbf{x}}^{(L)}]_{c^*} - y \right) \quad \text{subject to } y = \max_{t \neq c^*} [\hat{\mathbf{x}}^{(L)}]_t \wedge \text{MILP constraints},$$

where we have introduced a new variable y into the objective function.

Express the equality constraint

$$y = \max(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{K-1})$$

using only linear and integer constraints. To simplify notation, here $\mathbf{x}_k \in \mathbb{R}$ denotes the logit corresponding to the k -th incorrect class, and \mathbf{l}_k and \mathbf{u}_k its corresponding lower and upper bound.

Hint: You might want to introduce binary variables to indicate which logit is the maximum.

We first define $u_{max} := \max_k \mathbf{u}_k$, i.e. the largest upper bound.

Now we introduce the following constraints:

$$y \leq \mathbf{x}_k + (1 - b_k)(u_{max} - \mathbf{l}_k) \quad \forall 1 \leq k \leq K - 1 \quad (1)$$

$$y \geq \mathbf{x}_k \quad \forall 1 \leq k \leq K - 1 \quad (2)$$

$$\mathbf{b}_k \in \{0, 1\} \quad \forall 1 \leq k \leq K - 1 \quad (3)$$

$$\sum_{k=1}^{K-1} \mathbf{b}_k = 1 \quad (4)$$

The last constraint (4) simply ensures that only one element in \mathbf{b} is 1 and all others are zero.

The only valid assignment of \mathbf{b} is to have $\mathbf{b}_k = 1$ for the (unique) maximum value $\mathbf{x}_k = \max(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{K-1})$. To see this, consider the case that $\mathbf{b}_k = 1$ but \mathbf{x}_k is not the maximum value. Then, (1) resolves to $y \leq \mathbf{x}_k$. However, for the maximum value $\mathbf{x}_{max} > \mathbf{x}_k$ we have from (2) $y \geq \mathbf{x}_{max}$, leads to a contradiction.

Consider the case $\mathbf{b}_k = 1$ and the corresponding value \mathbf{x}_k is indeed the (unique) maximum. (1) and (2) imply that $y = \mathbf{x}_k$. The remaining values \mathbf{b}_i are zero, and in this case we need to show that (1) and (2) are never binding, regardless of the values \mathbf{x}_i . (2) is not binding since \mathbf{x}_i is not the maximum value. (1) is not binding because we have that $\mathbf{x}_i + u_{max} - \mathbf{l}_i \geq u_{max} \geq y$.

Problem 3: On slide 15 of the robustness chapter, we have defined an optimization problem for untargted attacks, i.e. we aim to have the sample $\hat{\mathbf{x}}$ classified as **any** class other than the correct one:

$$\min_{\hat{\mathbf{x}}} \mathcal{D}(\mathbf{x}, \hat{\mathbf{x}}) + \lambda \cdot L(\hat{\mathbf{x}}, y)$$

The loss function is defined as:

$$L(\hat{\mathbf{x}}, y) = \left[Z(\hat{\mathbf{x}})_y - \max_{i \neq y} Z(\hat{\mathbf{x}})_i \right]_+,$$

where $[\mathbf{x}]_+$ is shorthand for $\max(\mathbf{x}, 0)$ and $Z(\mathbf{x})_i = \log f(\mathbf{x})_i$ (i.e. log probability of class i). Here, $L(\hat{\mathbf{x}}, y)$ is positive if $\hat{\mathbf{x}}$ is classified correctly and 0 otherwise.

Provide an alternative loss function to turn this attack into a targeted attack, i.e. we aim to have the sample \mathbf{x} classified as a *specific* target class t .

$$L(\hat{\mathbf{x}}, t) = \left[\max_{i \neq t} Z(\hat{\mathbf{x}})_i - Z(\hat{\mathbf{x}})_t \right]_+$$

This loss is positive if $\hat{\mathbf{x}}$ is classified as a class that is **not** t , and is zero otherwise.

Problem 4: Recall from slide 41 the MILP constraints expressing the ReLU activation function:

$$\begin{aligned} y_i &\leq x_i - l_i(1 - a_i), \\ y_i &\leq u_i \cdot a_i, \\ y_i &\geq x_i, \\ y_i &\geq 0, \\ a_i &\in \{0, 1\}, \end{aligned}$$

where $u_i, l_i \in \mathbb{R}$ are upper and lower bounds on the value of the ReLU input x_i .

Show that – for an unstable unit (i.e. $u_i > 0 \wedge l_i < 0$) – a continuous relaxation on a leads to the convex relaxation constraints on slide 54. That is, replacing the constraint $a_i \in \{0, 1\}$ with $a_i \in [0, 1]$ yields

$$(u_i - l_i)y_i - u_i x_i \leq -u_i l_i.$$

We can combine the first two constraints on slide 41:

$$\begin{aligned} y_i &\leq x_i - l_i(1 - a_i) \\ y_i &\leq u_i \cdot a_i \end{aligned}$$

by expressing them as

$$y_i \leq \min(x_i - l_i(1 - a_i), u_i \cdot a_i).$$

Note that we are free to choose any value for a_i between 0 and 1. We want to choose a_i so that it leads to the loosest-possible constraint on y_i , since this leads to the maximum ‘leeway’ to optimize the objective function. More formally,

$$y_i \leq \max_{a_i} \min(x_i - l_i(1 - a_i), u_i \cdot a_i)$$

Further note that the two terms in the $\min(\cdot, \cdot)$ are two linear functions in a_i . Since $l_i < 0$, the first term is a function with negative slope in a_i . Since $u_i > 0$, the second term in the $\min(\cdot, \cdot)$ is a function with positive slope in a_i .

Consequently, the function $\min(x_i - l_i(1 - a_i), u_i \cdot a_i)$ is maximal at the intersection of the two linear functions. Solving for a_i we get:

$$\begin{aligned} x_i - l_i(1 - a_i) &= a_i u_i \\ \Leftrightarrow a_i &= \frac{x_i - l_i}{u_i - l_i} \end{aligned}$$

Plugging the expression of a_i into one of the original constraints, e.g. $y_i \leq a_i \cdot u_i$ we get:

$$y_i \leq \frac{x_i - l_i}{u_i - l_i} u_i$$

$$\Leftrightarrow y_i(u_i - l_i) - u_i x_i \leq -u_i l_i,$$

and therefore we have recovered the constraint of the convex relaxation.

Problem 5: Convex relaxations of non-linearities are not limited to ReLU. For this exercise, we consider the ReLU6 non-linearity

$$\text{ReLU6}(x) = \min(\max(0, x), 6),$$

which is used in MobileNet models performing low-precision computations on mobile devices.

Given input bounds l and u with $l \leq x \leq u$, provide a set of linear constraints corresponding to the convex hull of $\{(x \quad \text{ReLU6}(x))^T \mid l \leq x \leq u\}$.

Hint: You have to make a case distinction over different ranges of l and u .

To facilitate our discussion, we first rewrite the non-linearity as

$$\text{ReLU6}(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } 0 \leq x < 6 \\ 6 & \text{if } x \geq 6 \end{cases} \quad (5)$$

We can distinguish six different cases.

If $u < 0$ (and thus also $l < 0$), the non-linearity is always inactive (first case of Eq. 5). Thus, the convex hull is characterized by

$$y = 0.$$

If $l \geq 0$ and $u \leq 6$, we are always in the second case of Eq. 5 and thus

$$y = x.$$

If $l \geq 6$ (and thus also $u \geq 6$), the non-linearity is always saturated (third case of Eq. 5). In this case, we have

$$y = 6.$$

If $l < 0$ and $0 \leq u < 6$, ReLU6 behaves like an unstable ReLU unit. As discussed in lecture 4, the convex hull is a triangle spanned between the vertices $(l \quad 0)^T$, $(0 \quad 0)^T$ and $(u \quad u)^T$. It can thus be characterized by three linear constraints corresponding to its faces:

$$y \geq 0$$

$$y \geq x$$

$$y \leq \frac{u}{u-l}(x-l).$$

If $0 \leq l \leq 6$ and $u > 6$, the convex hull is – similar to the previous case – a triangle spanned between vertices $(l \ l)^T$, $(6 \ 6)^T$ and $(u \ 6)^T$. It can also be characterized by three linear constraints corresponding to its faces:

$$\begin{aligned} y &\leq x \\ y &\leq 6 \\ y &\geq l + \frac{6-l}{u-l}(x-l), \end{aligned}$$

with the last constraint corresponding to the line connecting $(l \ l)^T$ and $(u \ 6)^T$.

If $l < 0$ and $u > 6$, the convex hull has vertices $(l \ 0)^T$, $(0 \ 0)^T$, $(6 \ 6)^T$ and $(u \ 6)^T$. It can be characterized by the following four linear constraints corresponding to its faces:

$$\begin{aligned} y &\geq 0 \\ y &\leq 6 \\ y &\leq \frac{6}{6-l}(x-l) \\ y &\geq \frac{6}{u}x, \end{aligned}$$

with the third constraint corresponding to the line connecting $(l \ 0)^T$ and $(6 \ 6)^T$ and the fourth constraint corresponding to the line connecting $(0 \ 0)^T$ and $(u \ 6)^T$.

Randomized smoothing

Problem 6: (*) In the previous exercise we investigated the adversarial robustness of linear classifiers

$$f(\mathbf{x}) = \mathbb{I}[\mathbf{w}^T \mathbf{x} + b > 0]$$

with weight vector $\mathbf{w} \in \mathbb{R}^d$ and bias $b \in \mathbb{R}$, mapping samples from \mathbb{R}^d to binary labels $\{0, 1\}$.

Given such a linear classifier f , we can define the randomly smoothed classifier $g : \mathbb{R}^d \mapsto \{0, 1\}$ with

$$g(\mathbf{x}) = \operatorname{argmax}_{c \in \{0, 1\}} g_c(\mathbf{x})$$

and

$$g_c(\mathbf{x}) = \Pr_{\boldsymbol{\epsilon}}(f(\mathbf{x} + \boldsymbol{\epsilon}) = c) = \begin{cases} \Pr_{\boldsymbol{\epsilon}}(\mathbf{w}^T(\mathbf{x} + \boldsymbol{\epsilon}) + b \leq 0) & \text{if } c = 0 \\ \Pr_{\boldsymbol{\epsilon}}(\mathbf{w}^T(\mathbf{x} + \boldsymbol{\epsilon}) + b > 0) & \text{else} \end{cases},$$

where $\boldsymbol{\epsilon} \in \mathbb{R}^d$ is a random variable.

For this exercise, we assume that $\boldsymbol{\epsilon}$ follows an isotropic normal distribution, i.e. $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$ with elementwise standard deviation $\sigma \in \mathbb{R}_+$.

As discussed in the lecture, evaluating randomly smoothed classifier is typically not tractable and requires sampling. This is however not the case for our simple linear classifier.

Given input $\mathbf{x} \in \mathbb{R}^d$, weights $\mathbf{w} \in \mathbb{R}^d$ and bias $b \in \mathbb{R}$, show that $g_0(\mathbf{x}) = \Phi_{0,1} \left(-\frac{\mathbf{w}^T \mathbf{x}}{\sigma \|\mathbf{w}\|_2} - \frac{b}{\sigma \|\mathbf{w}\|_2} \right)$, where $\Phi_{0,1} : \mathbb{R} \mapsto [0, 1]$ is the cumulative distribution of the standard normal distribution $\mathcal{N}(0, 1)$.

Hint: $\Pr_{\epsilon} (\mathbf{w}^T(\mathbf{x} + \epsilon) + b \leq 0)$ can alternatively be written as:

$$\int_{\mathbb{R}^d} \mathbb{I} [\mathbf{w}^T(\mathbf{x} + \epsilon) + b \leq 0] \mathcal{N}(\epsilon \mid \mathbf{0}, \sigma^2 \mathbf{I}) d\epsilon.$$

We begin by deriving the expression for $g_0(\mathbf{x}) = \Pr_{\epsilon} (\mathbf{w}^T(\mathbf{x} + \epsilon) + b \leq 0)$. The smoothed output could be calculated by integrating over all possible values of ϵ :

$$g_0(\mathbf{x}) = \int_{\mathbb{R}^d} \mathbb{I} [\mathbf{w}^T(\mathbf{x} + \epsilon) + b \leq 0] \mathcal{N}(\epsilon \mid \mathbf{0}, \sigma^2 \mathbf{I}) d\epsilon. \quad (6)$$

We can however simplify our derivations by noticing that the value of the indicator function only depends on the value of the scalar random variable

$$z = \mathbf{w}^T(\mathbf{x} + \epsilon) + b.$$

Any affine transformation $\mathbf{A}\mathbf{y} + \mathbf{c}$ of a multivariate normal random variable $\mathbf{y} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ is also a multivariate normal random variable following distribution $\mathcal{N}(\mathbf{A}\boldsymbol{\mu} + \mathbf{c}, \mathbf{A}\boldsymbol{\Sigma}\mathbf{A}^T)$. Thus, the scalar random variable Z , which is the result of an affine transformation of ϵ follows a univariate normal distribution:

$$z \sim \mathcal{N}(\mu, \alpha)$$

with

$$\begin{aligned} \mu &= \mathbf{w}^T \mathbf{0} + \mathbf{w}^T \mathbf{x} + b = \mathbf{w}^T \mathbf{x} + b \\ \alpha &= \sqrt{\mathbf{w}^T \sigma^2 \mathbf{I} \mathbf{w}} = \sigma \sqrt{\mathbf{w}^T \mathbf{w}} = \sigma \|\mathbf{w}\|_2. \end{aligned}$$

Note that we take the square root since we parameterize the univariate distribution based on its standard deviation, not its variance.

Rather than integrating over the vector ϵ , we can instead integrate over the scalar z :

$$g_0(\mathbf{x}) = \int_{-\infty}^{\infty} \mathbb{I} [z \leq 0] \mathcal{N}(z \mid \mu, \alpha) dz = \int_{-\infty}^0 \mathcal{N}(z \mid \mu, \alpha) dz \quad (7)$$

From the right-hand side of Eq 7, it becomes clear that we are simply evaluating the cumulative distribution function $\Phi_{\mu, \alpha}$ of $\mathcal{N}(\mu, \alpha)$, i.e.

$$g_0(\mathbf{x}) = \Phi_{\mu, \alpha}(0).$$

The distribution $\mathcal{N}(\mu, \alpha)$ is a standard normal distribution $\mathcal{N}(0, 1)$ that is translated by μ and scaled by a factor α . Thus, we can alternatively write

$$g_0(\mathbf{x}) = \Phi_{0,1}((0 - \mu)/\alpha) = \Phi_{0,1} \left(-\frac{\mathbf{w}^T \mathbf{x}}{\sigma \|\mathbf{w}\|_2} - \frac{b}{\sigma \|\mathbf{w}\|_2} \right),$$

where the last equality simply follows from the definition of μ and α .

Randomized smoothing for discrete data

For the sake of simplicity, we consider a slightly different setup than in the lecture. In this exercise, we assume no knowledge about $f_\theta(\mathbf{x})$ respectively $g(\mathbf{x})_c$ (usually we would estimate a lower bound of $g(\mathbf{x})_c$ via Monte Carlo sampling, but here we do not).

We use the same sparsity-aware randomization scheme $\phi(\mathbf{x})$ as in the lecture:

$$g(\mathbf{x})_c = \mathcal{P}(f(\phi(\mathbf{x})) = c) = \sum_{\tilde{\mathbf{x}} \text{ s.t. } f(\tilde{\mathbf{x}})=c} \prod_{i=1}^{n^2} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i) \quad (8)$$

with

$$\mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i) = \begin{cases} p_d^{\mathbf{x}_i} p_a^{1-\mathbf{x}_i} & \tilde{\mathbf{x}}_i = 1 - \mathbf{x}_i \\ (1 - p_d)^{\mathbf{x}_i} (1 - p_a)^{1-\mathbf{x}_i} & \tilde{\mathbf{x}}_i = \mathbf{x}_i \end{cases} \quad (9)$$

and the number of nodes n . For an illustration we refer to Slide 15 “Smoothed Classifier for Discrete Data”

Problem 7: (*) Given an arbitrary graph \mathbf{x} , and a perturbed one \mathbf{x}' where \mathbf{x}' differs from \mathbf{x} in exactly one edge. What is the worst-case base classifier $h^*(\mathbf{x})$? In this context, we refer to the worst-case base classifier $h^*(\mathbf{x})$ as the classifier that has the largest drop in classification confidence between $g(\mathbf{x})_c$ and $g(\mathbf{x}')_c$. Or in other words, $h^*(\mathbf{x})$ results in the most instable smooth classifier if we switch a single edge. This motivates the importance of analyzing robustness for graph neural networks (or other models with discrete input data).

The classifier with the *largest drop in classification accuracy between $g(\mathbf{x})_c$ and $g(\mathbf{x}')_c$* can be formalized as a minimization problem $h^*(\mathbf{x}) = \arg \min_{h(\mathbf{x}) \in \mathcal{H}} g(\mathbf{x}')_c - g(\mathbf{x})_c$. In the following we consider a random order of edges and hence we may assume w.l.o.g. that all edges are identical but the last edge. Hence, from (8) it follows:

$$\begin{aligned} \min_{h(\mathbf{x}) \in \mathcal{H}} g(\mathbf{x}')_c - g(\mathbf{x})_c &= \min_{h(\mathbf{x}) \in \mathcal{H}} \left(\sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}})=c} \prod_{i=1}^{n^2} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}'_i) \right) - \left(\sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}})=c} \prod_{i=1}^{n^2} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i) \right) \\ &= \min_{h(\mathbf{x}) \in \mathcal{H}} \sum_{\substack{\tilde{\mathbf{x}} \text{ s.t.} \\ h(\tilde{\mathbf{x}})=c}} \left[\left(\prod_{i=1}^{n^2-1} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}'_i) \right) \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}'_{n^2}) - \left(\prod_{i=1}^{n^2-1} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i) \right) \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}_{n^2}) \right] \\ &= \min_{h(\mathbf{x}) \in \mathcal{H}} \sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}})=c} \left(\prod_{i=1}^{n^2-1} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i) \right) \underbrace{(\mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}'_{n^2}) - \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}_{n^2}))}_{\Delta_{\tilde{\mathbf{x}}}} \end{aligned}$$

$\Delta_{\tilde{\mathbf{x}}} = \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}'_{n^2}) - \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}_{n^2})$ resolves to two cases (each case occurs 50% of the time): (1) $1 - (p_a + p_d)$ and (2) $p_a + p_d - 1$. To minimize $g(\mathbf{x}')_c - g(\mathbf{x})_c$ we now choose $h^*(\mathbf{x})$ to predict c for all cases where $\Delta_{\tilde{\mathbf{x}}} < 0$ (assuming $p_a + p_d \neq 1$). Hence, $\Delta = \Delta_{\tilde{\mathbf{x}}}$ for $\tilde{\mathbf{x}}$ s.t. $h(\tilde{\mathbf{x}}) = c$.

We conclude the worst-case base classifier $h^*(\mathbf{x})$ exactly classifies exactly 50% of the random graphs $\tilde{\mathbf{x}}$ with c (note that in the general case $g(\mathbf{x})_c \neq 1/2$). In the case where one edge is removed from \mathbf{x}' (relatively to \mathbf{x}) and $p_a + p_d < 1$, the worst case base classifier $h^*(\mathbf{x})$ predicts c for all graphs where this edge is not missing (e.g. $h^*(\mathbf{x}) = c$ and $h^*(\tilde{\mathbf{x}}) \neq c$).

Problem 8: (*) How many of the possible graphs $\tilde{\mathbf{x}}$ does the worst-case base classifier assign the label c (see Problem 7)? To be more specific, we are looking for a term reflecting the absolute number and not a ratio?

Since we have n^2 edges there are 2^{n^2} possible adjacency matrices (each adjacency matrix represents one graph). Since we predict 50% with class c , we have a total of $2^{n^2}/2 = 2^{n^2-1}$ graphs resulting in c .

This clearly shows that enumerating all possible $\tilde{\mathbf{x}}$ is infeasible also for very small graphs.

Problem 9: What is $g(\mathbf{x}')_c$, $g(\mathbf{x})_c$, and $g(\mathbf{x}')_c - g(\mathbf{x})_c$ for the worst-case base classifier $h^*(\mathbf{x})$ (see Problem 1)? Please derive the equations (given $p_a + p_d < 1$). Subsequently, we would like to know the precise values for $p_a = 0.001$ and $p_d = 0.1$.

Since $p_a + p_d < 1$ we conclude that $\Delta = p_a + p_d - 1$.

$$\begin{aligned}
 \min_{h(\mathbf{x}) \in \mathcal{H}} g(\mathbf{x}')_c - g(\mathbf{x})_c &= \min_{h(\mathbf{x}) \in \mathcal{H}} \sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}})=c} \left(\prod_{i=1}^{n^2-1} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i) \right) \underbrace{(\mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}'_{n^2}) - \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}_{n^2}))}_{\Delta} \\
 &= \min_{h(\mathbf{x}) \in \mathcal{H}} \underbrace{\Delta \sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}})=c} \prod_{i=1}^{n^2-1} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i)}_{=1} \\
 &= p_a + p_d - 1
 \end{aligned}$$

Please note that $\sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}})=c} \prod_{i=1}^{n^2-1} \mathcal{P}(\tilde{\mathbf{x}}_i | \mathbf{x}_i)$ can be understood as a sum over the entire sample space of a product of $(n^2 - 1)$ Bernoulli random variables (i.e. sum over all possible combinations). Due to the basic laws of probability it must sum up to one.

Using $\Delta = \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}'_{n^2}) - \mathcal{P}(\tilde{\mathbf{x}}_{n^2} | \mathbf{x}_{n^2})$ s.t. $h^*(\tilde{\mathbf{x}}) = c$, we can easily go back and forth between $g(\mathbf{x}')_c$, $g(\mathbf{x})_c$, and $g(\mathbf{x}')_c - g(\mathbf{x})_c$. Consequently, the worst-case base classifier, with the given flip probabilities $p_a = 0.001$ and $p_d = 0.1$, has the following probabilities:

- $g(\mathbf{x}')_c = p_a = 0.001$
- $g(\mathbf{x})_c = 1 - p_d = 0.9$
- $g(\mathbf{x}')_c - g(\mathbf{x})_c = p_a + p_d - 1 = -0.899$

Please acknowledge that a smooth classifier might predict the right class c with high probability $g(\mathbf{x})_c = 1 - p_d = 0.9$, but flipping a single edge can result in $g(\mathbf{x}')_c = p_a = 0.001$. Hence, the probability of the smooth classifier drops by around 90%.