Chair for Computer Vision and Artificial Intelligence
Department of Informatics
Technical University of Munich

Esolution
Place student sticker here

Note:
- During the attendance check a sticker containing a unique code will be put on this exam.
- This code contains a unique number that associates this exam with your registration number.
- This number is printed both next to the code and to the signature field in the attendance check list.

Introduction to Deep Learning

Exam:	IN2346 / Endterm	**Date:**	Thursday 8th August, 2019
Examiner:	Prof. Dr. Leal-Taixé, Prof. Dr. Nießner	**Time:**	08:00 – 09:30

	P 1	P 2	P 3	P 4	P 5	P 6
I						

Working instructions

- This exam consists of **20 pages** with a total of **6 problems**.
  Please make sure now that you received a complete copy of the exam.
- The total amount of achievable credits in this exam is 90 credits.
- Detaching pages from the exam is prohibited.
- Allowed resources:
  - none
- Do not write with red or green colors nor use pencils.
- Physically turn off all electronic devices, put them into your bag and close the bag. This includes calculators.

Left room from _____ to _____ / Early submission at _____

Problem 1 **Multiple Choice (18 credits)**

Mark your answer clearly by a cross in the corresponding box. Multiple correct answers per question possible.

a) Your network is overfitting. What are good ways to approach this problem?

- [ ] Increase the size of the validation set
- [x] Increase the size of the training set
- [x] Reduce your model capacity
- [ ] Reduce learning rate and continue training

b) A sigmoid layer

- [ ] has a learnable parameter.
- [ ] cannot be used during backpropagation.
- [x] is continuous and differentiable everywhere.
- [x] maps to values between -1 and 1.

c) Training error does not decrease. What could be a reason?

- [x] Too much regularization.
- [x] Too many weights in your network.
- [x] Bad initialization.
- [x] Learning rate is too high.

d) How many network parameters are in ResNet-152?

- [ ] 1,337,337.
- [x] 60,344,232.
- [ ] more than a billion.
- [ ] 152.

e) What is the correct order of operations for an optimization with gradient descent?

a Update the network weights to minimize the loss.
b Calculate the difference between the predicted and target value.
c Iteratively repeat the procedure until convergence.
d Compute a forward pass.
e Initialize the neural network weights.

- [ ] bcdea
- [ ] ebadc
- [ ] eadbc
- [x] edbac

e d b a c

f) Dropout

- [ ] has trouble with tanh activations.
- [x] is an efficient way for regularization.
- [x] can be seen as an ensemble of networks.
- [ ] makes your network train faster.

g) Consider a simple convolutional neural network with a single convolutional layer. Which of the following statements is true about this network?

- [ ] All input nodes are connected to all output nodes.
- [x] It is scale invariant.
- [x] It is translation invariant.
- [x] It is rotation invariant.

h) You are building a model to predict the presence (labeled 1) or absence(labeled 0) of a tumor in a brain scan. The goal is to ultimately deploy the model to help doctors in hospitals. Which of these two metrics would you choose to use?

- [x] Recall = $\frac{\text{True positive examples}}{\text{Total positive examples}}$.
- [x] Precision = $\frac{\text{True positive examples}}{\text{Total predicted positive examples}}$.
- [x] Average Precision = $\frac{\text{True positive examples + True negative examples}}{\text{Total examples}}$.

i) Why you would want use 1×1 convolutions? (check all that apply)

- [ ] Predict binary class probabilities.
- [x] Collapse number of channels.
- [x] Learn more complex functions by introducing additional non-linearities.
- [x] To enforce a fixed size output.

25 min

Problem 2 Short Questions (24 credits)

0 1 2

a) You are training a neural network with 15 fully-connected layers with a *tanh* nonlinearity. Explain the behavior of the gradient of the non-linearity with respect to very large positive inputs.

Tanh:

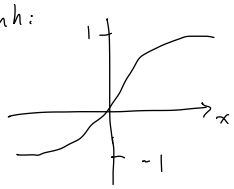

If the input of Activation layer is $\to +\infty$ or $-\infty$ then the tanH is equal to ± 1 Almost flat for large value

That means when doing gradient via Backpropagation, the gradient will equal to 0 be almost

0 1 2

b) Why might this be a problem for training neural networks? Name and explain this phenomenon.

This situation called Gradient ~~stopped~~ / Gradien Vanishing
Saturated

the weight of NN ~~won't~~ update because of the gradient is ~~equal~~ to 0
very slow

0 1 2

c) In modern architectures, another type of non-linearity is commonly used. Draw and name this non-linearity (1p) and explain why it helps solve the problem mentioned in the previous two questions (1p).

ReLU

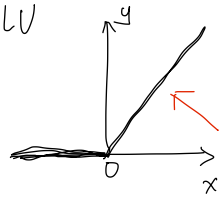

Linear
not saturates for large input

when $x \to -\infty$, output of Action laye is equal to 0
and when $x \to \infty$, output of — is equal to x.

then it won cause any gradient problem
gradient will not equal to 0

0 1 2

d) Why do we often refer to *L*2-regularization as "weight decay"? Derive a mathematical expression that includes the weights *W*, the learning rate η, and the *L*2-regularization hyperparameter λ to explain your point.

$W_{i+1} = (1-\lambda) W_i - \alpha \nabla L$

Weight update with Loss Function L incl weight decay

$W = W - \eta \nabla_w \left(L + \frac{1}{2} \lambda \Sigma_i W_i^2 \right)$

$W = W(1 - \eta\lambda) - \eta \nabla_w L$

$\eta\lambda << 1$

So the W is pushed toward to zero in each iteration

$W^t = W - \eta \nabla \left[L(w) + \frac{1}{2} \lambda \|w\|_2^2 \right]$

$= W - \eta \nabla L(w) - \eta\lambda w$

$= (1 - \eta\lambda) w - \eta \nabla L(w)$

$\eta\lambda << 1$ Therefore with weight update W will be force to equal to 0

e) You are solving the binary classification task of classifying images as cars vs. persons. You design a CNN with a single output neuron. Let the output of this neuron be z. The final output of your network, $\hat{y}$ is given by:

$$\hat{y} = \sigma\left(ReLU(z)\right)$$

You classify all inputs with a final value $\hat{y} \geq 0.5$ as car images. What problem are you going to encounter?

0 1

For sigmoid

Relu

because the output of ReLU is $\max(0, z)$
And z is $> 0 \Rightarrow \sigma \geq 0.5$
So All images will be classified as car,
if set $\hat{y} \geq 0.5$ as car images

f) Suppose you initialize your weights **w** with uniform random distribution $U(-\alpha, \alpha)$. The output **s** for given input vector **x** is given by

$$s_i = \sum_{j=0}^{n} w_{ij} \cdot x_j,$$

0 1 2

where n is the number of input values.
Assume that the input data **x** and weights are independent and identically distributed. How do you have to choose α such that the variance of the input data and the output is identical, hence $\mathrm{Var}(s) = \mathrm{Var}(x)$.
Hint: For two statistically independent variables X and Y holds:

$$\mathrm{Var}(X \cdot Y) = \left[\mathrm{E}(X)\right]^2 \mathrm{Var}(Y) + \left[\mathrm{E}(Y)\right]^2 \mathrm{Var}(X) + \mathrm{Var}(X)\mathrm{Var}(Y)$$

Furthermore the PDF of an uniform distribution $U(a, b)$ is

$$f(x) = \begin{cases} \frac{1}{b-a} & \text{for } x \in [a, b] \\ 0 & \text{otherwise.} \end{cases}$$

The variance of a continuous distribution is calculated as

$$\mathrm{Var}(X) = \int_{\mathbb{R}} x^2 f(x)\, dx - \mu^2,$$

where μ is the expected value of X.

$\mathrm{Var}(s) = \mathrm{Var}(\sum s_i) = \mathrm{Var}\left(\sum_{i=1}^{n}\sum_{j=0}^{n} w_{ij} x_j\right)$
$= \mathrm{Var}\left(\sum_{i=1}^{n} WX\right)$
$= n\, \mathrm{Var}(W)\, \mathrm{Var}(x)$
$= \mathrm{Var}(x)$
$\mathrm{Var}(W) = \frac{1}{n} = \frac{1}{3}\alpha^2$
$\alpha = \sqrt{\frac{3}{n}}$

$\mathrm{Var}(s) = \mathrm{Var}(Wx) = \left[E(W)\right]^2 \mathrm{Var}(X) + \left[E(X)\right]^2 \mathrm{Var}(W) + \mathrm{Var}(W)\mathrm{Var}(X)$

$\mathrm{Var}(s_i) = \mathrm{Var}\left(\sum_{j=0}^{n} w_{ij} \cdot x_j\right)$
$= n\, \mathrm{var}(W \cdot X)$

$W = \sim U(a, b) = \begin{cases} \frac{1}{2\alpha} & x \in [-\alpha, \alpha) \\ 0 \end{cases}$

$E(X) = \int x \cdot f(x)\, dx$
$= \int_a^b x \cdot \frac{1}{b-a}\, dx$
$= \frac{x^2}{2(b-a)} = \frac{b^2 - a^2}{2(b-a)} = \frac{b+a}{2} = \frac{-\alpha+\alpha}{2} = 0$

$\mathrm{Var}(W) = \int_{\mathbb{R}} \frac{\alpha}{2}\, d\alpha - \mu^2$
$= \frac{1}{4}\alpha^2 - \mu^2$

$\mathrm{Var}(X) = \int x^2 f(x) dx - \mu^2$
$= \int_a^b \frac{x^3}{3(b-a)} - \left(\frac{a+b}{2}\right)^2$
$= \frac{b^3 - a^3}{3(b-a)} - \frac{(a+b)^2}{4}$
$= \frac{(b-a)^2}{12} = \frac{1}{3}\alpha^2$

$\int x^2 f(x) dx = \frac{1}{3}(b-a)(b^2+ab+a^2)$
$\frac{(b+a)^2}{4}$

$\int_a^b \frac{1}{2} x^2 \Big|_a^b$
$\frac{1}{b-a}$
$\frac{1}{2}(b^2 - a^2)$
$= \frac{1}{2}(b+a)$

$x \in [a, b)$
$U \sim [-\alpha, \alpha)$
$E(x) = 0$

$\mathrm{Var}(x) = \frac{4b^2+4ab+4a^2}{12} - \frac{3b^2+3ab+3a^2}{12}$
$= \frac{(b-a)^2}{12}$

$\mathrm{Var}(W) = \frac{\alpha^2}{3}$
$n \cdot \frac{\alpha^2}{3} = 1$
$\alpha = \sqrt{\frac{3}{n}}$

g) Consider 2 different models for image classification of the MNIST data set. The models are: (i) a 3 layer perceptron, (ii) LeNet.
Which of the two models is more robust to translation of the digits in the images? Give a short explanation why.

0 1 2

LeNet, convolution layers

0 1 2

h) Consider the following one-dimensional data points with classes $\{0, 1\}$. Sketch a linear (0.5p) and logistic (0.5p) regression into the figures. Which model is more suitable for this task (1p)?

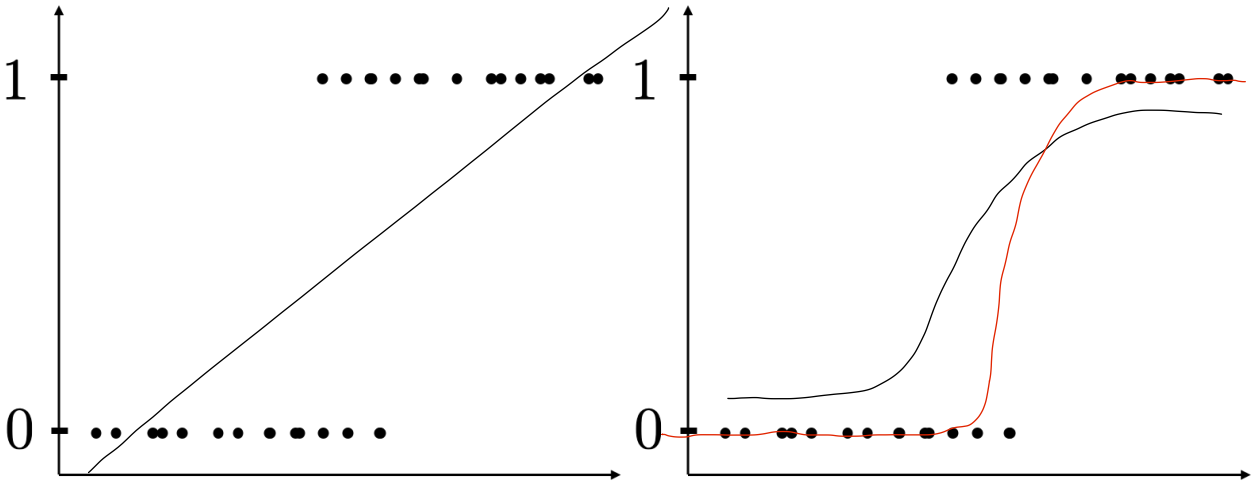

Plot linear regression (left) and logistic regression (right).

Regression because the label is $\{0, 1\}$

0 1 2 3

i) What is the mean and standard deviation of Xavier initialization? What changes to this initialization would you propose when used with ReLU non-linearities?

$Var(x) = \frac{1}{x}$ Mean = 0

Use kaming -Inalizatio

$Var(x) = \frac{2}{x}$

NN prefer Cat because imbalance between classes

0 1 2

j) You have 4000 cat and 100 dog images and want to train a neural network on these images to do binary classification. What problems do you foresee with this dataset distribution? Name two possible solutions.

The minibatch may have different distribution, $\Rightarrow$ the NN will cause Training error not decrease

collect more dog picture

reweight data loden

Data Augmentation for dog picture / Dropout

0 1 2

k) Why is initializing all the weights of a fully connected layer to the same value problematic during training?

All neuron have same weights $\Rightarrow$ the gradient via backpropagate on each neuron are same operation $\Rightarrow$ the gradient are same

$\Rightarrow$ ~~learn nothing~~

Node learn the same thing during back propagation

l) What is the difference between dropout for convolutional layers compared to dropout for fully connected layers? Explain both behaviours.

0 1 2

Dropout ann in FCL means use only half (p=0.5) neurons to train model

Conv drop feature map random / Fully connected drop weights at random

30 min

Problem 3 Optimization (12 credits)

a) Explain the concept behind RMSProp optimization. How does it help converging faster? (0 / 1 / 2)

RMSProp use second-order Momentum, which called Look-ahead-momentum. It will predict the location which the optimum will arrive and jump to that location and do some correction

b) Which SGD variation uses first and second momentum? (0 / 1)

Adam

c) Why is it common to use a learning rate decay? (0 / 1)

We want fast convergence when it's far away from global optimum so that we can speed up the training process. At meanwhile, we want slow convergence while it's near to the global optimum so that we won't miss or jump over the global optimum point

d) What is a saddle point? What is the advantage/disadvantage of Stochastic Gradient Descent (SGD) in dealing with saddle points? (0 / 1 / 2)

The goal, the global optimum

Saddle, the gradient is zero neither local minima or local maxima

SGD use stochastic weight update, can help escape saddle point

Disadvantage: SGD will take long time to get the saddle point because of stochastic

e) Why would one want to use larger mini-batches in SGD? (0 / 1)

Large mini-batches means in every iteration, ~~SGD~~ will learn more than and SGD will make a better convergence

+ More efficient computation / better generalization

Make gradient less noisy

faster convergence

- sensitive to the weight Initialization

f) Why do we usually use small mini-batches in practice? (0 / 1)

We have too many data, it's very time and hardware costly when doing training on large-batches

Limited GPU Memory, faster compute, faster update

g) Your network's training curve diverges (assuming data loading is correct). Name one way to address the problem through hyperparameter ~~change~~. (0 / 1)

Use small learning rate

reduce

h) What is an epoch?

0
1

One Epoch means the optimation update te weights when using all mini batch through out all training data / mini batch in one Epoch

full run through entire training set.

i) When is SGD guaranteed to converge to a local minima (provide formula)?

0
1
2

$\sum_{i=1}^{\infty} \alpha_i = \infty$ and $\sum_{i=1}^{\infty} \alpha_i^2 < \infty$ $\alpha_n \geq 0 \quad \forall n \geq 0$

Problem 4 Convolutional Neural Networks and Advanced Architectures (12 credits)

In the following we assume that the input of our network is a $224 \times 224 \times 3$ color (RGB) image. The task is to perform image classification on 1000 classes. You design a network with the following structure [CONV - RELU] x 20 - FC - FC. That is, you place 20 consecutive convolutional layers (including non-linear activations), followed by two fully-connected layers. Each layer will have its own number of filters and kernel size.

0 1

a) The first 3 convolutional layers have each 5 filters with kernels of size 3×3, applied with stride 1 and no padding. How large is the receptive field of a feature after the 3 convolutional operations?

Input: $\frac{F_1 - N}{S} + 1 = F_2$ $F_1 = 1$

$\frac{F_2 - N}{S} + 1 = F_3$ $F_2 = 5$

$\frac{F_3 - N}{S} + 1 = 1$ $F_3 = 7$ 7×7

0 1 2

b) What are the dimensions of the feature map after the 3 convolutional operations from (a) ?

$\frac{F_1 - N}{S} + 1 = F_2$ $\frac{224 - 3}{1} + 1 = 222$ $218 \times 218 \times 5$

$\frac{222 - 3}{1} + 1 = 220$

$\frac{220 - 3}{1} + 1 = 218$

0 1 2

c) What are the dimensions of the weight tensor of the first convolutional layer? (1p) What does each dimension represent? (1p)

$3 \times 3 \times 5$ (width, height, channel)

$\times$ (3, 5, 3, 3) (filter size, input channel, output size)

0 1 2

d) After the 10th convolutional layer your feature map has size 100x100x224. You realize the next convolutional filter operation will involve too many multiplications that make your network training slow. However, the next layer requires identical *spatial size* of the feature map.
Propose a solution for this problem (1p) and demonstrate your solution with an example (1p).

Use Depth wise separable convolutional layer

Use 1x1 convolutional layer to reduce the channel

Use 1x1x256, 25 filter

then get 100 x 100 x 25

cheaper and faster computation for next layer

0 1 2

e) Your network is now trained for the task of image classification. You now want to use the trained weights of this network for the task of *image segmentation* for which you need a pixel-wise output. Which layers of your original network described above can you *not* reuse for the image segmentation task? (1p) Describe briefly how you would adapt the network for image segmentation given *any input image size*? (1p)

~~last~~ FC layer, which is softmax because they take fixed input.

fine adapt the input size to the origin network input size by another layer or pooling

converted FC → to Fully Convolution as up sampling

f) You decide to increase the number of layers substantially and therefore you switch to a ResNet architecture. Draw a ResNet block (1p). Describe all the operations inside the block (1pt). What is the advantage of using such a block in terms of training (1p)?

0 1 2 3

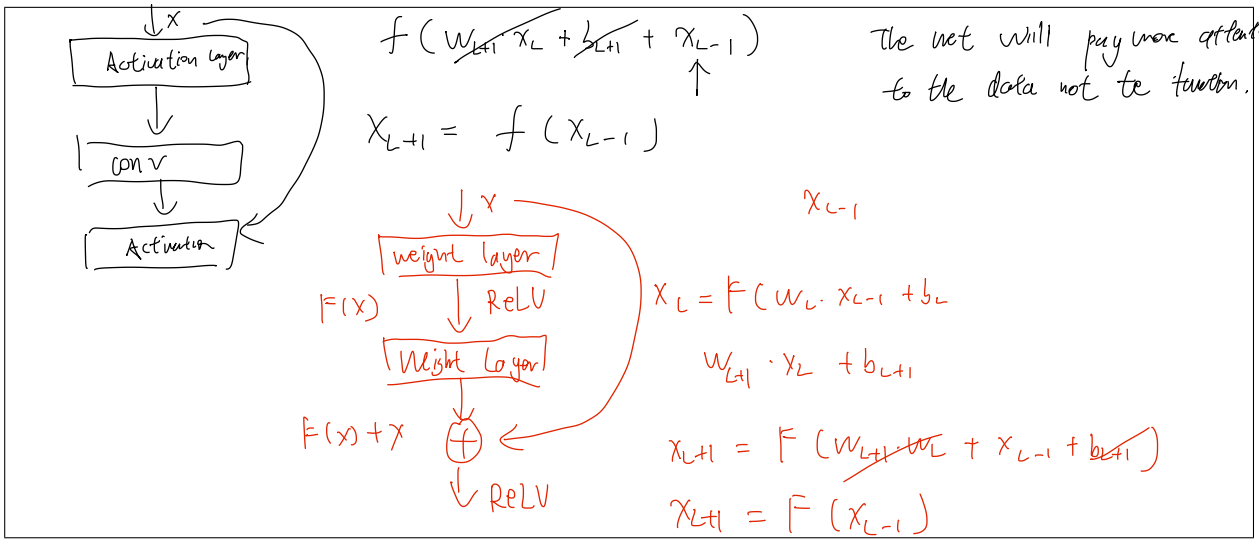

Sum the input and the output from last Activation layer

resolving gradient vanishing

highway for gradient / deeper network easier to train

Problem 5 Backpropagation and Convolutional Layers (12 credits)

Your friend is excited to try out those "Convolutional Layers" you were talking about from your lecture. However, he seems to have some issues and requests your help for some theoretical computations on a toy example.
Consider a neural network with a convolutional (without activation) and a max pooling layer. The convolutional layer has a single filter with kernel size $(1, 1)$, no bias, a stride of 1 and no padding. The filter weights are all initialized to a value of 1. The max pooling layer has a kernel size of $(2, 2)$ with stride 2, and 1 zero-padding.

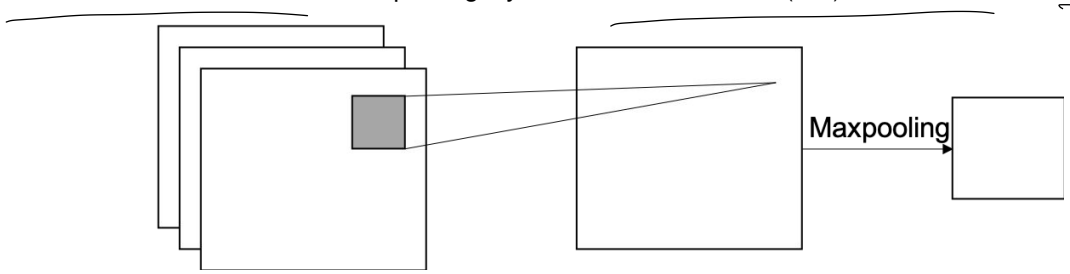

You are given the following input image of dimensions $(3, 2, 2)$:

$$x = \left(\begin{bmatrix} 1 & -0.5 \\ 2 & -2 \end{bmatrix}, \begin{bmatrix} -2 & 1 \\ -1.5 & 1 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix} \right)$$

0 1 2 a) Compute the forward pass of this input and write down your calculations.

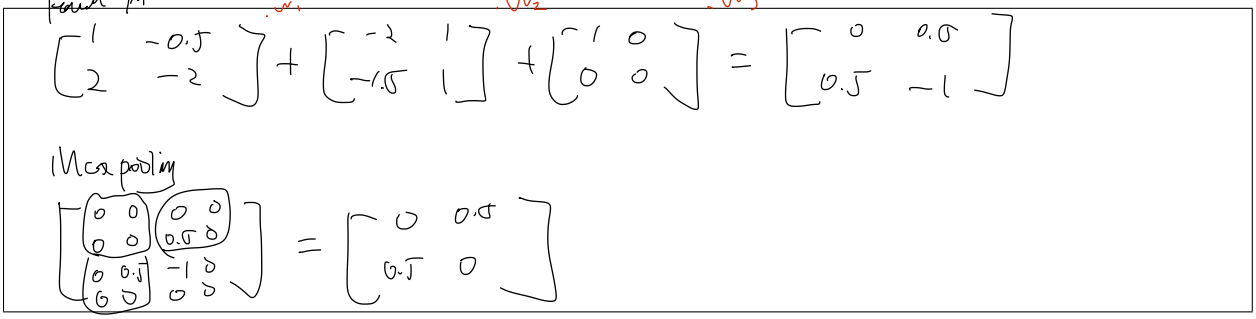

0 1 b) Consider the corresponding ground truth,

$$y = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$

Calculate the binary cross-entropy with respect to the natural logarithm by summing over all output pixels of the forward pass computed in (a). You may assume $\log(0) \approx -10^9$. (Write down the **equation** and keep the logarithm for the final result.)

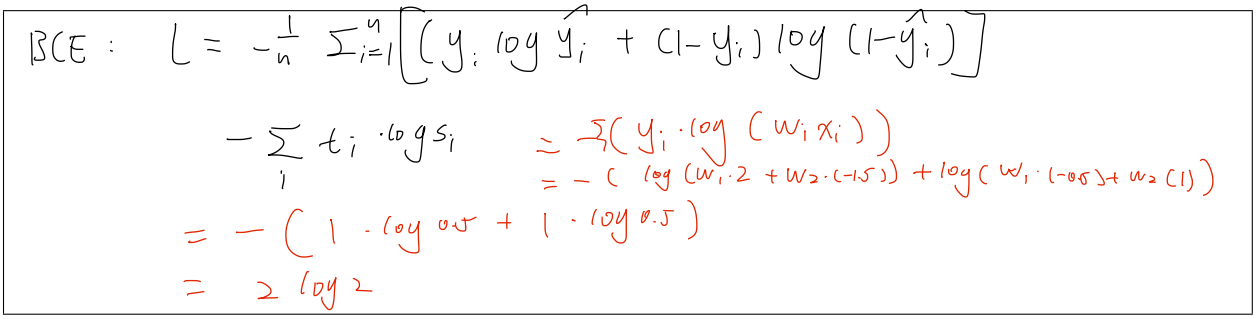

0 ½ c) You don't recall learning the formula for backpropagation through convolutional layers but those 1×1 convolutions seem suspicious. Write down the name of a common layer that is able to produce the same result as the convolutional layer used above.

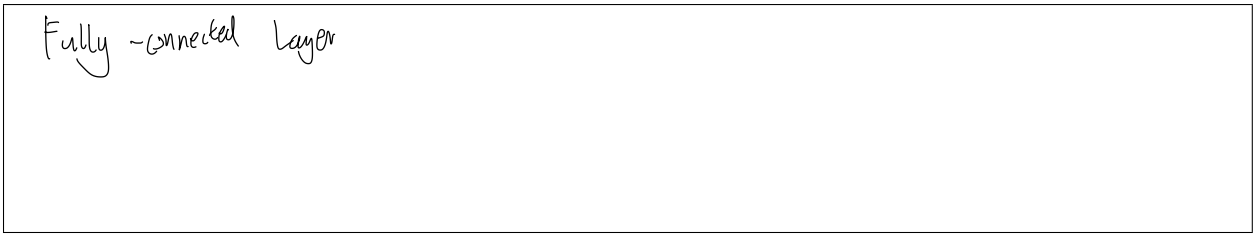

d) Update the kernel weights accordingly by using gradient descent with a learning rate of 1. (Write down your calculations!)

$$\frac{\partial BCE}{\partial w_1} = \frac{-\left(\log(w_1 \cdot 2 + w_2 \cdot (-1.5)) + \log(w_1 \cdot (-0.5) + w_2 \cdot 1)\right)}{\partial w_1}$$

$$= -\frac{2}{w_1 \cdot 2 + w_2 (-1.5)} - \frac{(-0.5)}{w_1 \cdot (-0.5) + w_2 \cdot 1} = -4 + 1 = -3$$

$$\frac{\partial BCE}{\partial w_2} = -\frac{(-1.5)}{w_1 \cdot 2 + w_2 (-1.5)} - \frac{1}{w_1 \cdot (-0.5) + w_2 (1)} = 3 - 2 = 1$$

$$w_1 = w_1 - \alpha \cdot \frac{\partial BCE}{\partial w_1} = 1 - 1 \cdot (-3) = 4$$

$$w_2 = w_2 - \alpha \frac{\partial BCE}{\partial w_2} = 1 - 1 \cdot 1 = 0$$

$(3, 1, 1, 1) = 3$ 个参量

$$\frac{\partial BCE}{\partial w_3} = 0$$

$$w_3 = w_3 - \alpha \frac{\partial BCE}{\partial w_3} = 1 - 1 \cdot 0 = 1$$

0 1 2

e) After helping your friend debugging, you want to showcase the power of convolutional layers. Deduce what kind of 3×3 convolutional filter was used to generate the output (right) of the grayscale image (left) and write down its 3×3 values.

Edge detection

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & 8 \\ -1 & -1 & -1 \end{bmatrix}$$

Vertical edge detector

Prewitt kernel

0 1

f) He finally introduces you to his real problem. He wants to find 3×3 black crosses in grayscale images, i.e., each pixel has a value between 0 (black) and 1 (white).

You notice that you can actually hand-craft such a filter. Write down the numerical values of a 3×3 filter that maximally highlights on the position of black crosses.

$$\begin{bmatrix} -1 & 1 & -1 \\ 1 & -1 & 1 \\ -1 & 1 & -1 \end{bmatrix}$$

Problem 6 Recurrent Neural Networks and LSTMs (12 credits)

a) Consider a vanilla RNN cell of the form $h_t = \tanh(V \cdot h_{t-1} + W \cdot x_t)$. The figure below shows the input sequence x_1, x_2, and x_3.

0
1
2

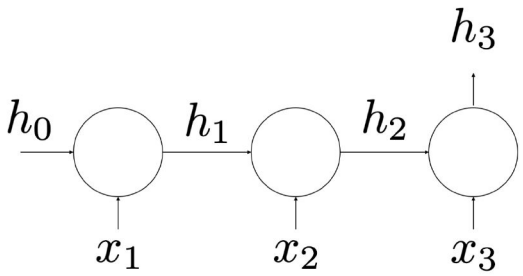

Given the dimensions $x_t \in \mathbb{R}^4$ and $h_t \in \mathbb{R}^{12}$, what is the number of parameters in the RNN cell? Neglect the bias parameter.

output $\rightarrow$ 12

$4 \times 12 + 12 \times 12 = 192$

b) If x_t is the 0 vector, then $h_t = h_{t-1}$. Discuss whether this statement is correct.

0
1
2

False

there has Weight V and Non-linearity tanh

so $ht \neq ht-1$

0
1
2
3

c) Now consider the following **one-dimensional** ReLU-RNN cell.

$$h_t = \text{ReLU}(V \cdot h_{t-1} + W \cdot x_t)$$

(Hidden state, input, and weights are scalars)
Calculate h_1, h_2 and h_3 where $V = 1$, $W = 2$, $h_0 = -3$, $x_1 = 1$, $x_2 = 2$ and $x_3 = 0$.

For $t = 1$

$V \cdot h_0 + W \cdot x_1 = 1 \cdot (-3) + 2 \cdot 1 = -1$

$\text{ReLU}(-1) = 0$

$h_1 = 0$

For $t = 2$

$V \cdot h_1 + W \cdot x_2 = 1 \cdot (0) + 2 \cdot 2 = 4$

$\text{ReLU}(4) = 4$

$h_2 = 4$

For $t = 2$

$V h_2 + W x_3 = 1 \cdot 4 + 2 \cdot 0 = 4$

$\text{ReLU}(4) = 4$

$h_3 = 4$

d) Calculate the derivatives $\frac{\partial h_3}{\partial V}$, $\frac{\partial h_3}{\partial W}$, and $\frac{\partial h_3}{\partial x_1}$ for the forward pass of the ReLU-RNN Cell of (c). Use that $\frac{\partial}{\partial x}\text{ReLU}(x)\Big|_{x=0} = 1$.

0 1 2 3

$h_t = \text{ReLU}(V \cdot h_{t-1} + W x_t) = \text{ReLU}(z_t)$

$$\frac{\partial h_3}{\partial V} = \frac{\partial \text{ReLU}(z_3)}{\partial z_3} \cdot \frac{\partial z_3}{\partial V} + \frac{\partial \text{ReLU}(z_2)}{\partial z_2} \cdot \frac{\partial z_2}{\partial V} + \frac{\partial \text{ReLU}(z_1)}{\partial z_1} \cdot \frac{\partial z_1}{\partial V}$$

$$= 1 \cdot h_2 + 1 \cdot h_1 + 0 \cdot h_0$$

$$= 4$$

$$\frac{\partial h_3}{\partial W} = \frac{\partial \text{ReLU}(z_3)}{\partial z_3} \cdot \frac{\partial z_3}{\partial W} + \frac{\partial \text{ReLU}(z_2)}{\partial z_2} \cdot \frac{\partial z_2}{\partial W} + \frac{\partial \text{ReLU}(z_1)}{\partial z_1} \cdot \frac{\partial z_1}{\partial W}$$

$$= 1 \cdot x_3 + 1 \cdot x_2 + 0 \cdot x_1$$

$$= 2$$

$$\frac{\partial h_3}{\partial x_1} = \frac{\partial \text{ReLU}(z_3)}{\partial z_3} \cdot \frac{\partial z_3}{\partial h_2} \cdot \frac{\partial h_2}{\partial z_2} \cdot \frac{\partial z_2}{\partial h_1} \cdot \frac{\partial h_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial x_1}$$

$$= 1 \cdot V \cdot 1 \cdot V \cdot 0 \cdot W = 0$$

0 1 2

e) A Long-Short Term Memory (LSTM) unit is defined as

output $o = g_1 = \sigma(W_1 \cdot x_t + U_1 \cdot h_{t-1})$,
forget $f = g_2 = \sigma(W_2 \cdot x_t + U_2 \cdot h_{t-1})$,
update $u = g_3 = \sigma(W_3 \cdot x_t + U_3 \cdot h_{t-1})$,

$$\tilde{c}_t = \tanh(W_c \cdot x_t + u_c \cdot h_{t-1}),$$
$$c_t = g_2 \circ c_{t-1} + g_3 \circ \tilde{c}_t,$$
$$h_t = g_1 \circ c_t,$$

where g_1, g_2, and g_3 are the gates of the LSTM cell.
1) Assign these gates correctly to the **forget** f, **update** u, and **output** o gates. (1p)
2) What does the value c_t represent in a LSTM? (1p) long-term memory

Cell state

Additional space for solutions–clearly mark the (sub)problem your answers are related to and strike out invalid solutions.